

CAI 4104/6108 – Machine Learning Engineering: ML Engineering (1)

Prof. Vincent Bindschaedler

Spring 2024

Administrivia

- Homework 0 is due **today** (by 11:59PM)
- **No class** on Friday (1/19)
 - ◆ Exercise 1 — I will **pre-record** it

Background & Notation

- Scalars: a, b, c (if domain not specified, assume real number)
- Vectors: \mathbf{a}, \mathbf{b}
 - ◆ $a^{(i)}$ or a_i is the i^{th} element of the vector \mathbf{a} ; e.g.: if $\mathbf{b} = [7, 3]$ then $b_1=7$ and $b_2=3$
 - ◆ $\|\mathbf{a}\|$ is the Euclidean norm of the vector \mathbf{a} , i.e.: $\|\mathbf{a}\| = [\sum_i (a_i)^2]^{1/2}$
 - ✿ More generally: $\|\mathbf{a}\|_p = (\sum_i |a_i|^p)^{1/p}$ (called L_p -norm)
 - ◆ Dot product: $\mathbf{a} \cdot \mathbf{b} = \sum_i a_i b_i$
 - ◆ \mathbf{a}^T is the transpose of \mathbf{a}
- Matrices: \mathbf{A}, \mathbf{B}
 - ◆ \mathbf{a}_i is the i^{th} row of \mathbf{A}
 - ◆ a_{ij} is the j^{th} element of the i^{th} row of \mathbf{A}
 - ◆ \mathbf{A}^T is the transpose of \mathbf{A}
- Sets: \mathcal{A}, \mathcal{B}
 - ◆ $[n] = \{1, 2, \dots, n\}$
 - ◆ Union: $\mathcal{A} \cup \mathcal{B}$; intersection $\mathcal{A} \cap \mathcal{B}$; set difference: $\mathcal{A} \setminus \mathcal{B}$; cardinality: $|\mathcal{A}|$

■ Functions: $f: \mathcal{X} \rightarrow \mathcal{Y}$

- ◆ \mathcal{X} : domain; \mathcal{Y} : co-domain, range, or image set
- ◆ Functions can have a **global minimum** (or maximum) and several **local minimum** (or maximum)
- ◆ Derivatives and gradients
 - ✿ $f'(x)$ denotes the **derivative** of $f(x)$ [e.g., if $f(x)=x^2$ then $f'(x)=2x$]
 - ✿ If f takes multiple inputs (e.g., $f(x,y)$ or $f(\mathbf{x})$), then the **gradient** of f , denoted ∇f is the **vector of partial derivatives**
 - ✿ For example for $f(x,y)$: $\nabla f = [\partial f / \partial x, \partial f / \partial y]$

■ Random variables: $X: \Omega \rightarrow \mathcal{R}$

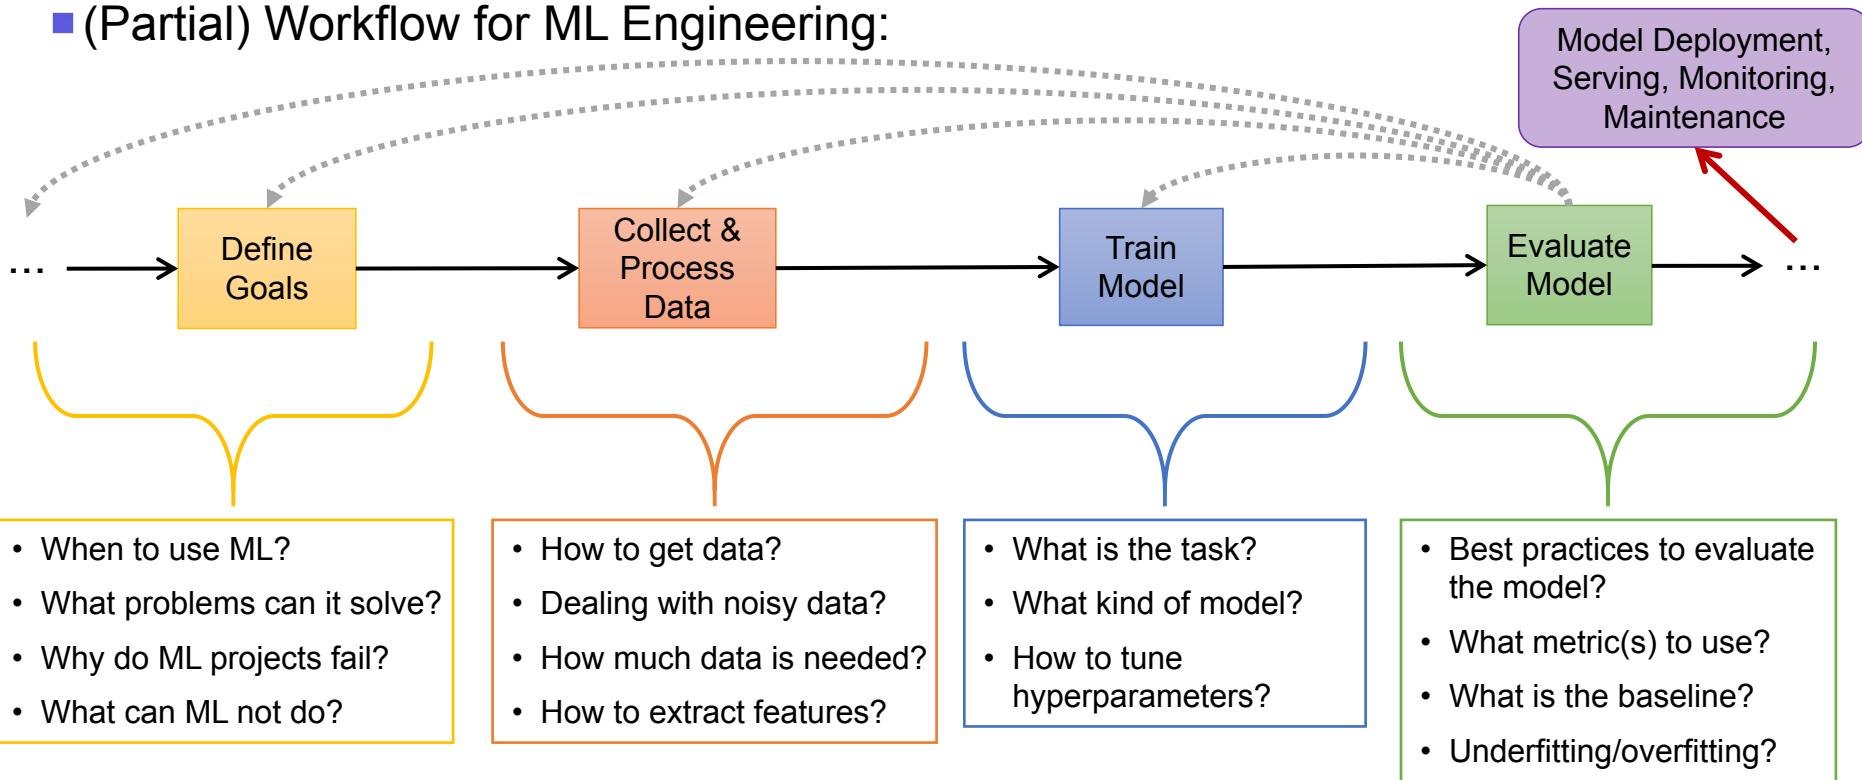
- ◆ X is a function from the sample space to some output set \mathcal{R}
- ◆ Discrete X : $\Pr\{X = x\}$ denotes the probability that X takes value $x \in \mathcal{R}$
 - ✿ The probability distribution of X is called the **probability mass function** (pmf)
- ◆ Continuous X : $\Pr\{X = x\}=0$ for any specific value $x \in \mathcal{R}$
 - ✿ We have a **probability density function** (pdf) denoted $p(x)$
- ◆ Expectation: $\mathbb{E}[X]$; Variance: $\text{Var}(X)$

■ Machine Learning

- ◆ Training dataset with n examples and m features: \mathbf{X} ($n \times m$ matrix)
 - ✧ \mathbf{x}_i represents the i^{th} example; $x_{i,j}$ or $x_i^{(j)}$ is the j^{th} feature value
- ◆ If supervised learning, then there is a corresponding set of labels \mathbf{y} ($n \times 1$ vector)
 - ✧ If doing classification, then the label of the i^{th} example y_i would be an integer (or encoded as one)
 - ✧ If doing regression, then y_i could be a real number
- ◆ Alternatively the training dataset may be denoted as: $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$
- ◆ Model: $\theta \in \Theta$
 - ✧ Prediction function $f_\theta(\mathbf{x})$ or $h_\theta(\mathbf{x})$
 - ✧ Loss function or cost function: L_θ or $J(\theta)$
 - For example: L_2 loss (aka “**squared error loss**”): $L = [y - f_\theta(\mathbf{x})]^2$

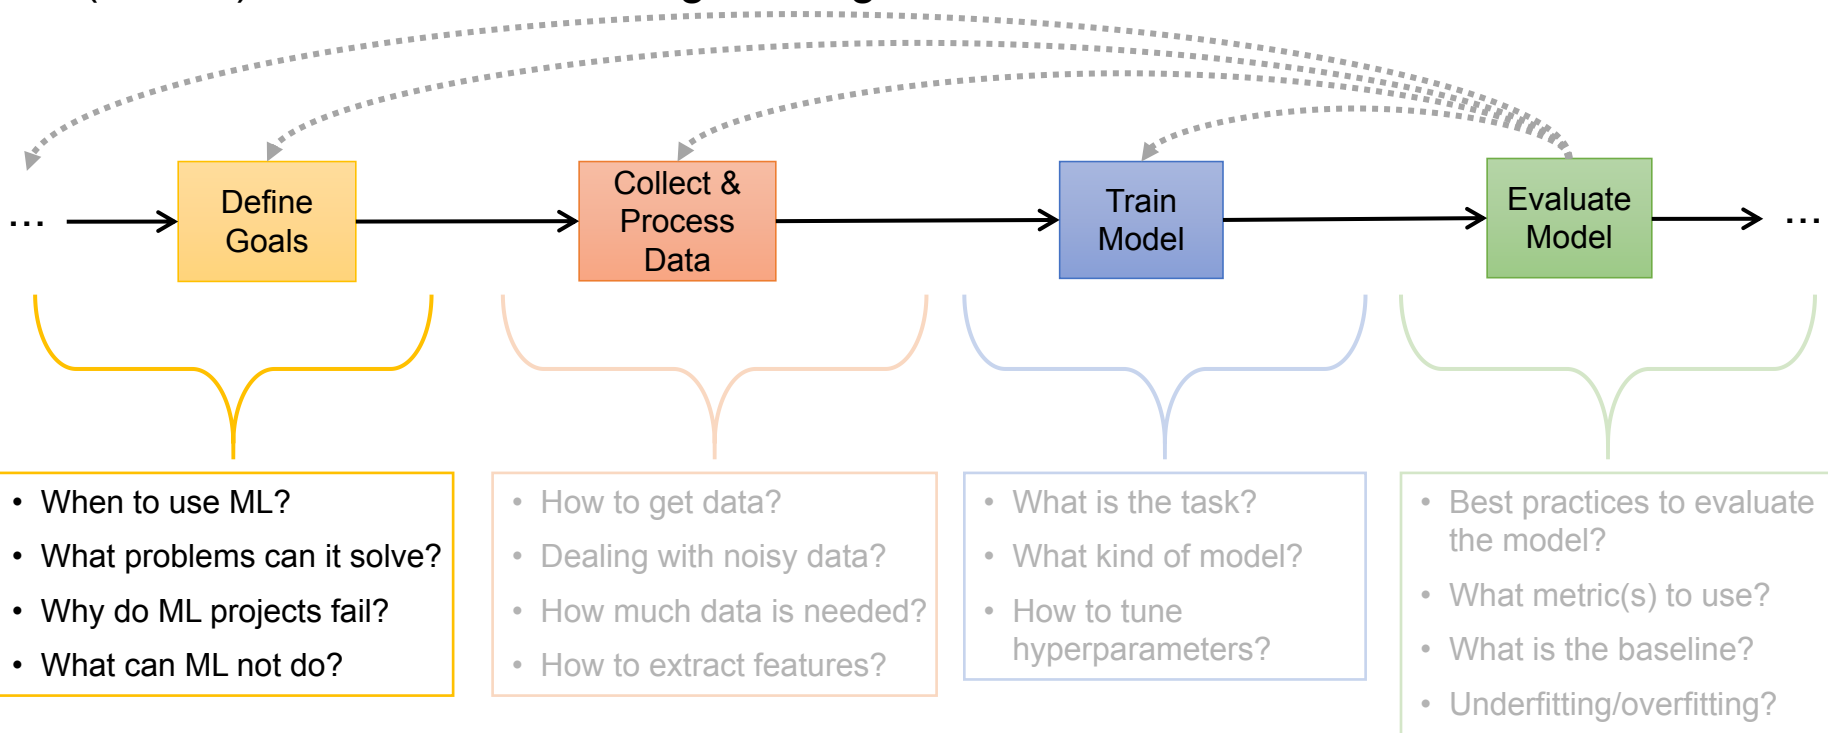
Machine Learning Engineering

■ (Partial) Workflow for ML Engineering:



Machine Learning Engineering

■ (Partial) Workflow for ML Engineering:



When Should We Use ML?

■ Consider using ML when:

- ◆ The problem is too complex for hardcoding rules
 - ✿ There could be too many rules, the rules may be unknown/hard to specify, or there could be too many parameters determine the rules
 - ✿ For example: what rules would you use for spam filtering?
- ◆ The problem deals with an unstudied/understudied phenomenon
 - ✿ Discovering unknown patterns in the data, maybe patterns difficult for humans to see
 - ✿ For example: predicting human behavior
- ◆ The problem calls for automating some decision/prediction
 - ✿ In order to reduce human workload or amount of data to look through
 - ✿ For example: anomaly detection in intrusion detection system
- ◆ The problem (or some aspect of it) is changing frequently
 - ✿ The data for a problem could be constantly changing (e.g., network traffic patterns)
 - ✿ We can (re)train model on new data, or we can use online/incremental learning

Many Reasons Not To Use ML

- Cannot get the (right) data or enough of it
- Problem does not require learning from data
- You can solve the problem in other ways
 - ◆ By developing new algorithms, using software development, etc.
- Cannot afford the cost of a mistake
- Judged **unethical/undesirable** to use ML
 - ◆ Should we use ML to determine prison/jail time?
- You need explanations, not just predictions
 - ◆ The problem requires **explaining some phenomon** (e.g., physics)
 - ✿ ML is **not** magic: it doesn't have deep insights about the natural world!
 - ✿ It's (just) a set of techniques and tools to make accurate predictions from data
 - ◆ *Note: there is active research on interpretable/explainable ML; but many models are black boxes!*
- Many others...

Why Do ML Projects Fail?

- According to VentureBeat AI, 87% of data science projects fail

- ◆ Source: <https://venturebeat.com/2019/07/19/why-do-87-of-data-science-projects-never-make-it-into-production/>

- Why do ML projects fail?

- ◆ Many reasons

- ✿ Not having people with the right expertise; insufficient computing infrastructure

- ◆ Bad data, or no data, or not enough data, or biased data

- ◆ Don't have the right data; or right data does not exist

- ◆ ML is the wrong approach

- ◆ Goal is technically infeasible

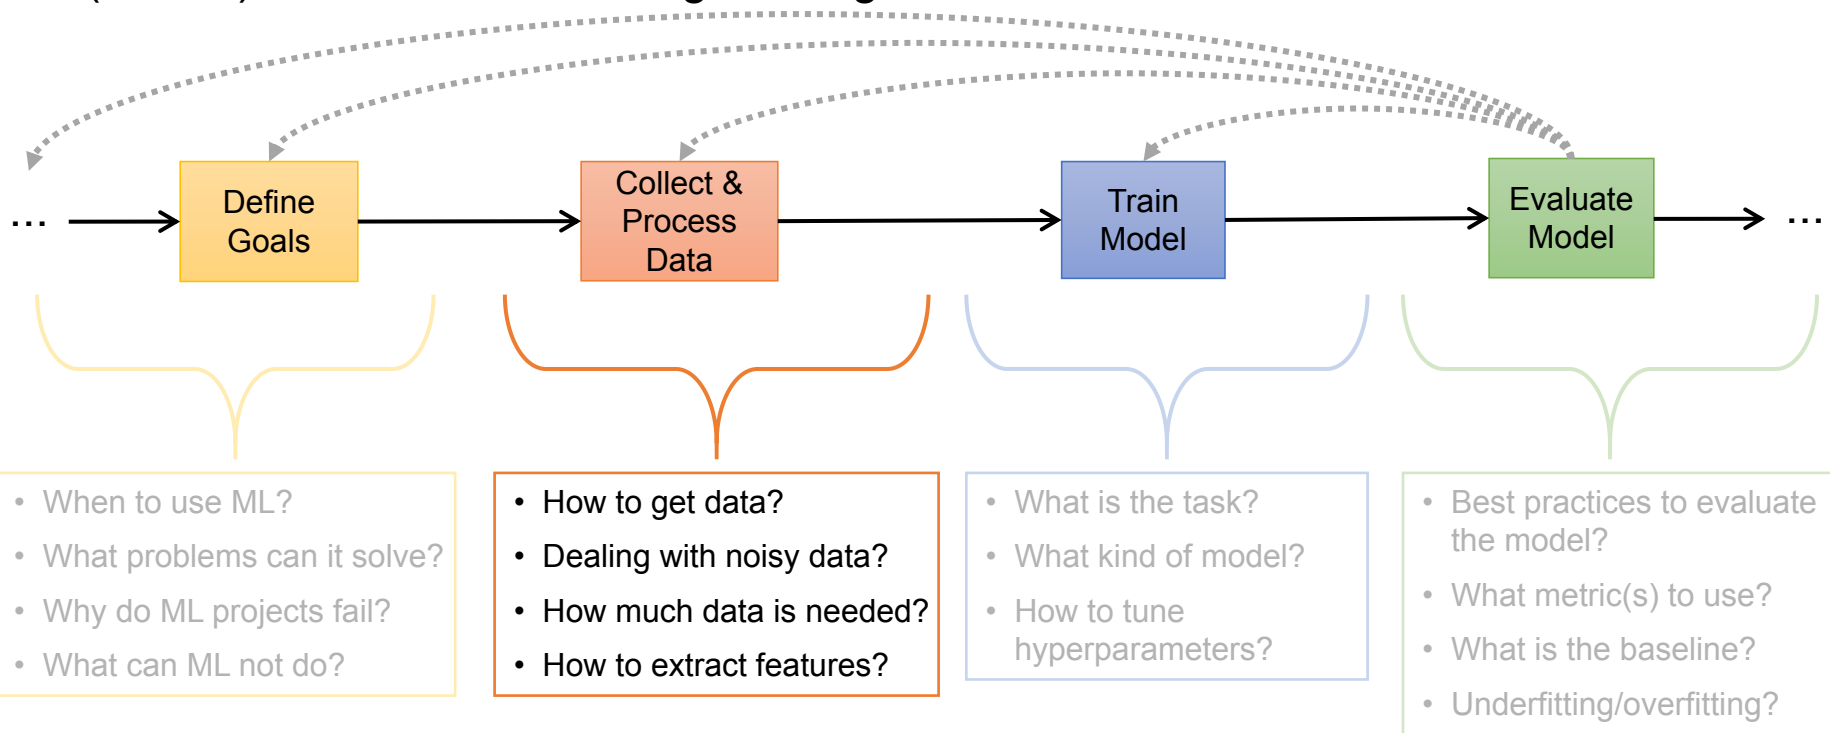
- ◆ Lack of metrics for success or bad metrics

- ✿ E.g., lack of baseline

- ◆ Many others...

Machine Learning Engineering

■ (Partial) Workflow for ML Engineering:




Collecting Data

■ Many questions:

- ◆ Where is the data coming from? Is it reliable?
- ◆ Does it need to be anonymized? (What can happen if we don't?)
- ◆ Is the data noisy/incomplete? Is it missing values?
- ◆ Is the data biased in some way?
 - ✧ Is the data expected to be similar to the data available after the system is deployed?

■ How much data do we need?

- ◆ More than you expect
- ◆ Rules of thumb:
 - ✧ More examples/instances than attributes/features
 - Some models (e.g., SVM) still work with few examples but lots of features
 - ✧ Many examples (e.g., 30+ or 50+) for each class
- ◆ How to answer this in practice? => **learning curves**
 - ✧ You'll be able to try this in assignment1



Dealing With Unclean Data

- The data is **noisy**
 - ◆ For example: some attribute values/features are wrong
 - ✿ Some learning algorithms are more sensitive than others
- The data has missing features values
 - ◆ Mitigation strategies
 - ✿ **Algorithmically**: use a learning algorithm that can deal with this (e.g., some Decision Trees implementations)
 - ✿ **Removal**: discard examples with missing values
 - ✿ (If feature is categorical/nominal) **Treat missing as value**: replace missing values with a “special value” (e.g., -1)
 - ✿ **Imputation**: use a data imputation technique (e.g., resampling from marginals, replace with average)
- The data contains duplicates
- The data is imbalanced
 - ◆ Weighting classes (if allowed by learning algorithm); oversampling; undersampling
- The data is incomplete or not representative?
 - ◆ Find better data!

Too Little Data? Too Much Data?

■ Data Augmentation

- ◆ Idea: **augment** the existing training data by using the training data!
 - ✿ Can have a **major** impact on model performance
 - ✿ But it won't help if you truly don't have enough data!
- ◆ For example: for images: rotation, scaling, crop, flip, etc.
- ◆ Data augmentation techniques are usually specific to the data type



■ Data Sampling

- ◆ If you have a lot of data, it may not be practical (and necessary) to use all of it
- ◆ Different strategies for sampling
 - ✿ **Random sampling**: e.g., pick a uniformly random subset of instances, pick each instance independently with some probability p
 - ✿ **Stratified sampling**: divide dataset into groups (**strata**), then randomly sample from each stratum
 - Note: this can be used to reduce bias (i.e., make the dataset more representative)

Data Augmentation?

- Data Augmentation can be thought of a **regularization** technique

Single sample

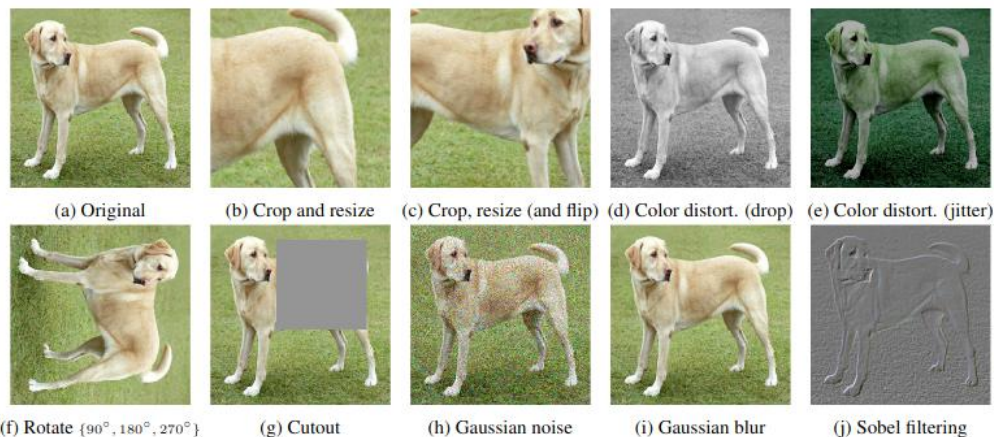
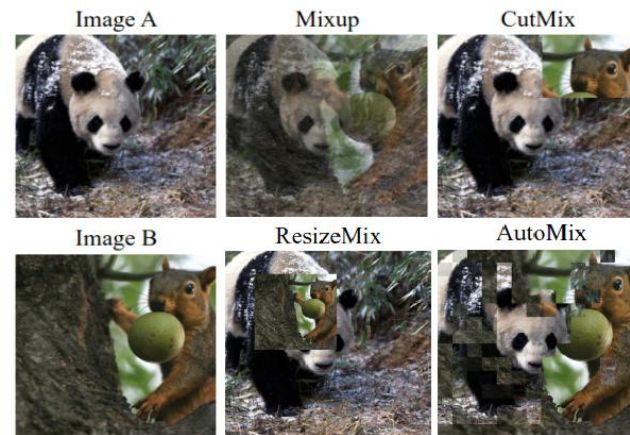


Figure 4. Illustrations of the studied data augmentation operators. Each augmentation can transform data stochastically with some internal parameters (e.g. rotation degree, noise level). Note that we *only* test these operators in ablation, the *augmentation policy used to train our models* only includes *random crop (with flip and resize)*, *color distortion*, and *Gaussian blur*. (Original image cc-by: Von.grzanka)

Source: Chen et al. "A simple framework for contrastive learning of visual representations." ICML 2020.

Multi-sample



Source: Li et al. "Openmixup: Open mixup toolbox and benchmark for visual representation learning." arXiv preprint arXiv:2209.04851.

■ Feature Engineering

- ◆ The process of transforming raw data into a dataset we can use
 - ✿ Time consuming process
 - ✿ Techniques are often specific to the data type you are dealing with (e.g., feature engineering for text)
- ◆ Goal: select a set of features that are highly predictive
 - ✿ But also not correlated among each other. **Why?**

■ Common Techniques

- ◆ For numerical features:
 - ✿ **Binning** (aka bucketing): aggregate range of numerical values into discrete bins
 - ✿ **Normalization**: remap values onto a range such as $[0,1]$ (**min-max normalization**) or $[-1,1]$
 - ✿ **Standardization** (aka z-score normalization): rescale feature values to follow a *standard normal distribution* (how? \Rightarrow subtract the mean and divide by the standard deviation)
- ◆ For categorical features:
 - ✿ **One-hot encoding**: turn a categorical feature with k possible values into a vector of k binary features
 - ✿ **Ordinal encoding**: if values are ordered/ranked, values can be encoded in order (e.g., using integers)

■ For numerical features:

- ◆ **Binning** (aka bucketing): aggregate range of numerical values into discrete bins
- ◆ **Normalization**: remap values onto a range such as $[0, 1]$ or $[-1, 1]$
- ◆ **Standardization** (aka z-score normalization): rescale feature values to follow a *standard normal distribution* (how? \Rightarrow subtract the mean and divide by the standard deviation)

■ For categorical features:

- ◆ **One-hot encoding**: turn a categorical feature with k possible values into a vector of k binary features with hamming weight 1
- ◆ **Ordinal encoding**: if values are ordered/ranked, values can be encoded in order (e.g., as integers)

■ Rules of thumb:

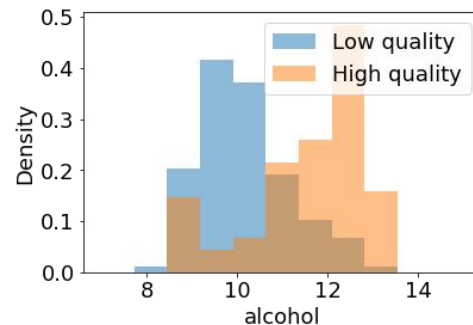
- ◆ If you have a mix of numerical and categorical features: use one-hot encoding / ordinal encoding
- ◆ Feature scaling (i.e., normalization & standardization) helps with most learning algorithms
 - ✱ For SVM (especially linear SVM), linear models (or if you are doing regularization): **you should rescale features!**

■ Preprocessing, data cleaning, and feature engineering

- ◆ Avoid manipulating the data in ad-hoc (i.e., non-reproducible ways) such as bash scripting, awk, manual editing, etc.
- ◆ For reproducibility, it is best to implement all the steps in a (Python) script
 - ✧ Ideally, you want to develop a **pipeline**

■ Visualizing data:

- ◆ Before getting too far into data cleaning or training a model, take a look at the (training) data!
 - ✧ Visualizing the data may yield insights into what features to use and what kind of model to use
 - ✧ You can also look at feature statistics: mean, standard deviation, min, max, mode, missing values, etc.
- ◆ Visualization examples:
 - ✧ Histograms,
 - ✧ Heatmaps, correlation plots,
 - ✧ Scatter plots, QQ-plots,
 - ✧ Class distribution plot per feature



Next Time

- Friday (1/19): Exercise 1

- ◆ *No Class - Pre-recorded*

- Upcoming:

- ◆ Homework 0 (due today by 11:59pm)