

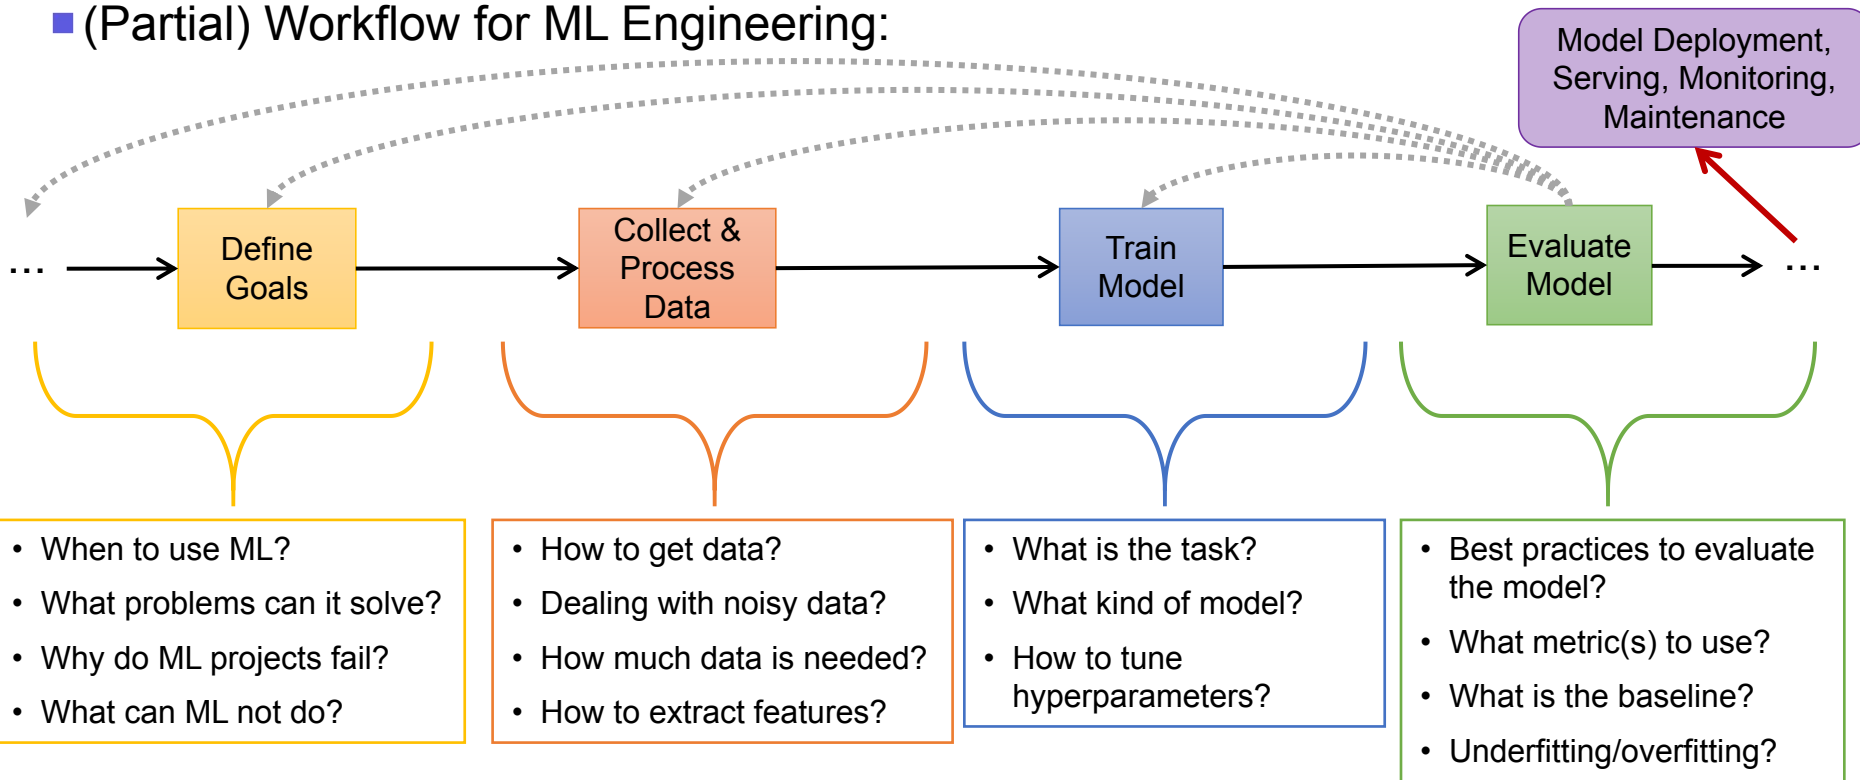
CAI 4104/6108 – Machine Learning Engineering: ML Engineering (3)

Prof. Vincent Bindschaedler

Spring 2024

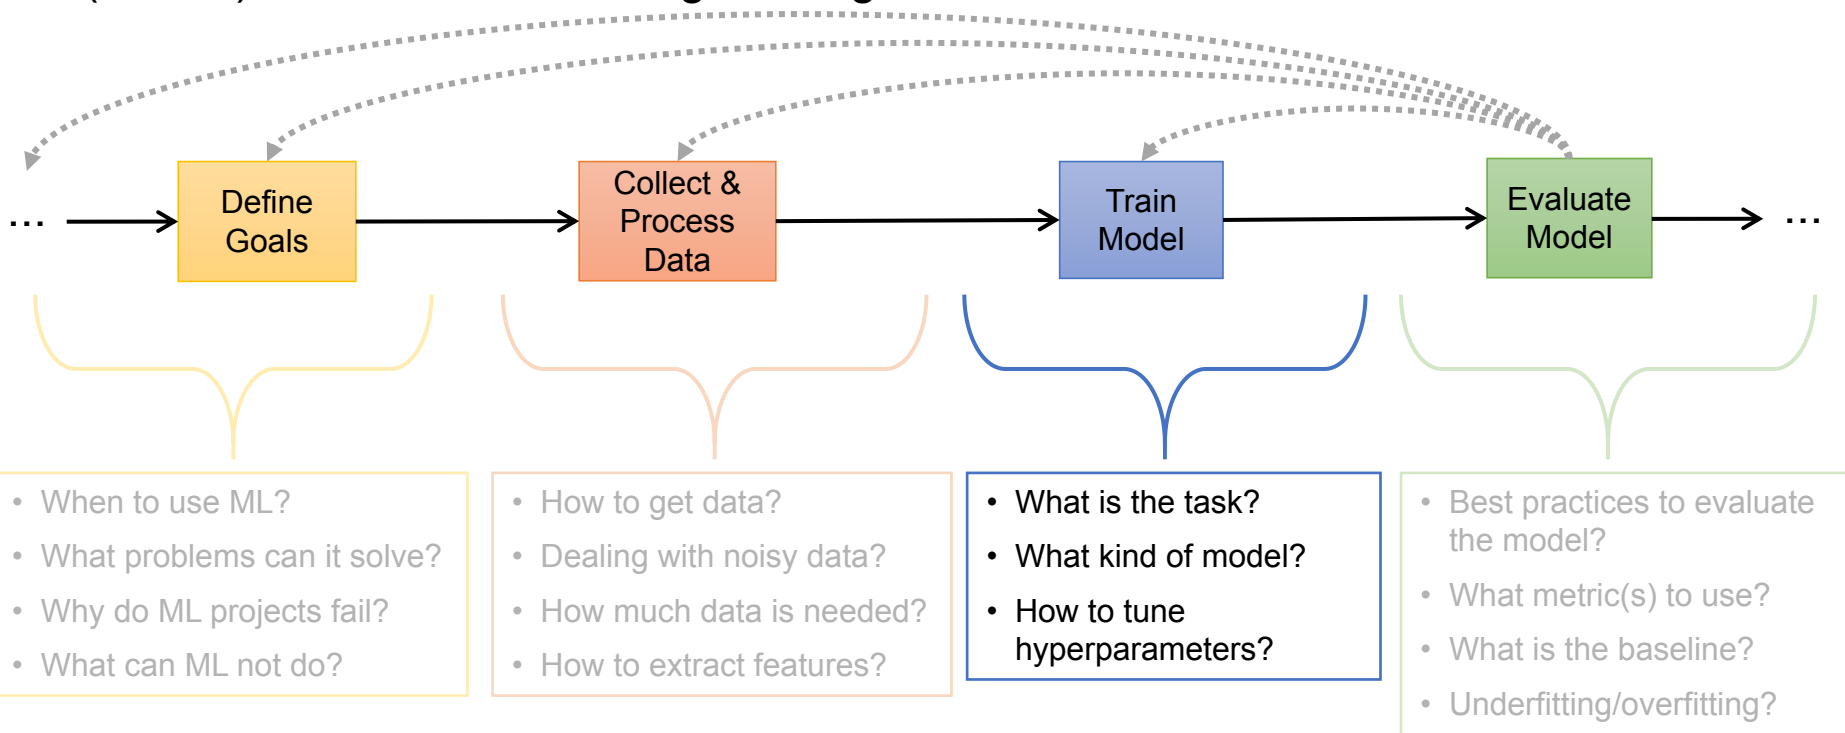
Reminder: Machine Learning Engineering

■ (Partial) Workflow for ML Engineering:



Machine Learning Engineering

■ (Partial) Workflow for ML Engineering:



Reminder: Types of Learning

■ Supervised Learning

- ◆ Learning from **labeled data** (i.e., each example or instance in the dataset has a corresponding label)
- ◆ Tasks: **classification** vs. **regression**

■ Unsupervised Learning

- ◆ Learning from **unlabeled data** (we must **discover patterns** in the data)
- ◆ Tasks: clustering (e.g., K-means), dimensionality reduction (e.g., t-SNE, PCA), etc.

■ Semi-supervised Learning

- ◆ Learning from **partially labeled data**

■ Reinforcement Learning

- ◆ There is an **agent** that can interact with its **environment** and perform **actions** and get **rewards**
- ◆ Learning a policy (i.e., a strategy for which actions to take) to get the most rewards over time


■ Transfer Learning

- ◆ Learning to **repurpose** an existing model for a new task

Reminder: Supervised Learning


■ Classification

- ◆ Task: predict the corresponding **label**
- ◆ Different types:
 - ✧ **Binary classification**: there are only two classes (0,1 ; +,-, etc.)
 - ✧ **Multiclass**: more than two classes
 - ✧ **Multi-label**: each instance can belong to more than one class (e.g., label all objects in a photo)
 - ✧ **One-class**: there is only one class, we want to distinguish it from everything else



■ Regression

- ◆ Task: predict the corresponding **value** (typically a real number) or **target**
 - ✧ E.g.: you want to predict a person's future income based on their high school GPA



■ Others:

- ◆ Sequence-to-sequence, similarity learning/metric learning, learning to rank, etc.

Reminder: Multiclass Classification

- **Multiclass classification** (aka *multinomial classification*)
 - ◆ There are $c > 2$ distinct classes: $y_i \in \{1, 2, \dots, c\}$ is the label of the i^{th} example
- Wait. How do we do this?
 - ◆ Recall the SVM formulation:
 - ✱ We want to find a **hyperplane** $w \cdot x - b = 0$, also we relabel the classes as +1 and -1. **Problem?**
- Some learning algorithms / models naturally support multiclass classification
 - ◆ E.g.: kNN, decision trees, neural networks
- For others, we can transform multiclass classification into a binary classification
 - ◆ **One-vs-rest** (OvR): Train c binary classifiers. f_i to classify class i versus not i
 - ✱ Predict: $y = \operatorname{argmax}_i f_i(x)$
 - ◆ **One-vs-one** (OvO): Train $c(c-1)/2$ binary classifiers. $f_{i,j}$ to classify class i versus class j
 - ✱ Predict: predict with all $c(c-1)/2$ and return the class that has the highest number of “votes”

What Is a Learning Algorithm?

■ To train a model we need:

- ◆ Some kind of objection function or criterion often called **loss function** (or **cost function**)
 - ✧ Some algorithms have a specific criterion (e.g., SVM — **hinge loss**)
 - ✧ Others do not have such a criterion (e.g., kNN)
- ◆ An optimization procedure to find a solution (i.e., parameter values)
 - ✧ E.g.: quadratic programming, gradient descent


■ We want “nice” loss functions:

- ◆ Ideally: **continuous**, **differentiable**, (strictly) **convex**, **smooth** loss functions
- ◆ Examples:
 - ✧ SVM is a convex problem
 - ✧ But (in general) the loss function for neural networks is **not** convex


Convex Sets & Convexity

■ Convex Sets:


Set \mathcal{X} is **convex** if for any $a, b \in \mathcal{X}$ the line $\lambda a + (1-\lambda)b \in \mathcal{X}$ for $\lambda \in [0,1]$



convex





not convex




convex

■ Convexity:


Function f on a convex set \mathcal{X} is **convex** if for any $x, y \in \mathcal{X}$ and $\lambda \in [0,1]$:
$$\lambda f(x) + (1-\lambda)f(y) \geq f(\lambda x + (1-\lambda)y)$$



strictly convex



convex



not convex

No Free Lunch Theorem

■ Famous ML result

◆ No Free Lunch (NFL) Theorem

✿ *David H. Wolpert. "The lack of a priori distinctions between learning algorithms." Neural computation 8.7 (1996).*

■ What does the theorem say?

◆ Informally: Given two learning algorithms A and B, there are just as many problem instances/datasets where A performs better than B as vice-versa (B performs better than A)

◆ In other words:

✿ There is **no** learning algorithm that is **guaranteed** to work well on our data a priori (i.e., before we try it)

■ Why?

◆ When we train a model (i.e., use a learning algorithm as opposed to another), we are making assumptions about how features are related to target/label

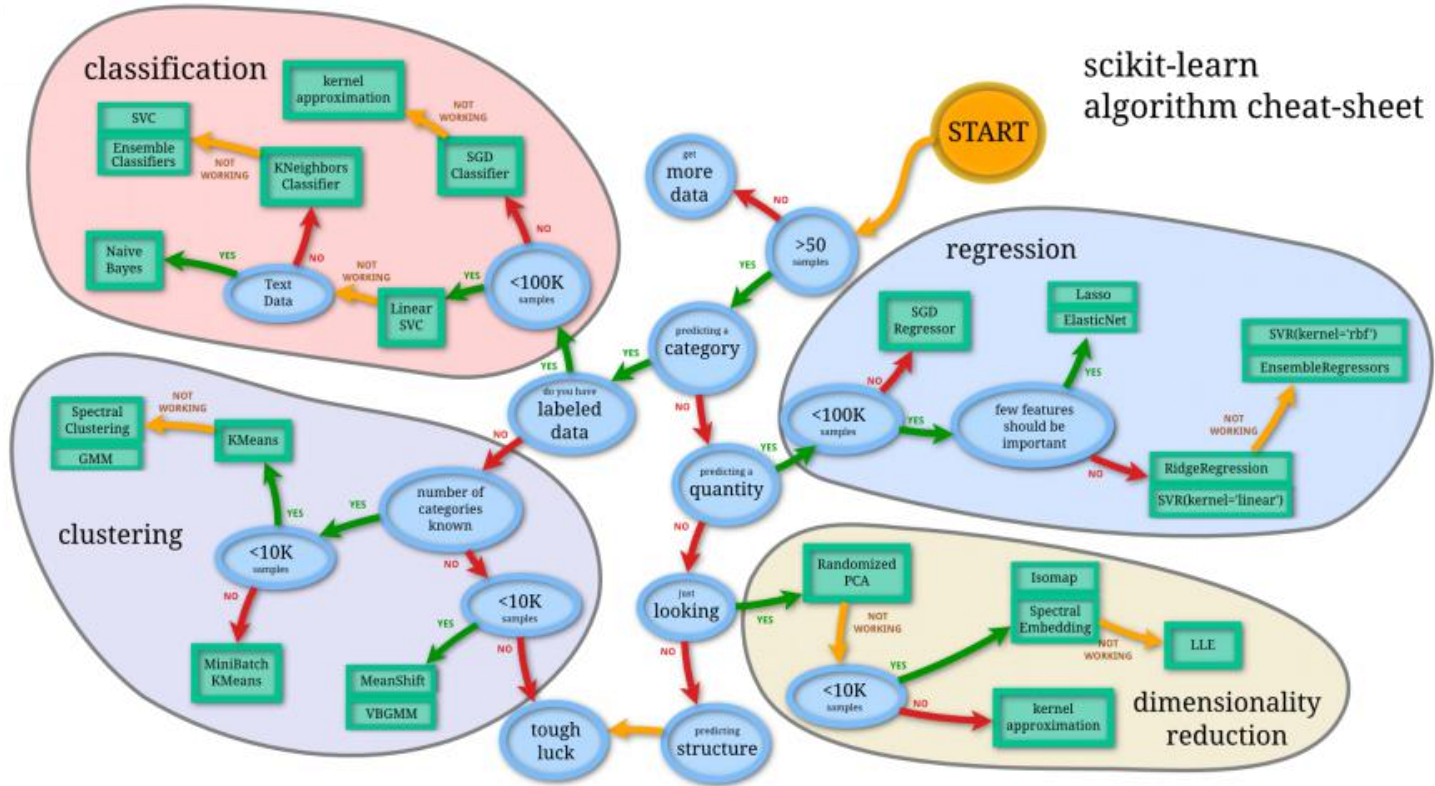
✿ E.g.: we assume two classes can be separated by a linear decision boundary based features (if we use linear SVM)

■ Consequences?

◆ The only way to know for sure what learning algorithm is best is to evaluate them all

◆ Or in practice: start with reasonable assumptions, then evaluate (only) a few algorithms

Model Selection



Source: scikit-learn.org

Parameters & Hyperparameters

■ Parameters

- ◆ Sometimes called “weights”
- ◆ Model parameters are determined by the training data (e.g., through some optimization procedure)

■ Hyperparameters

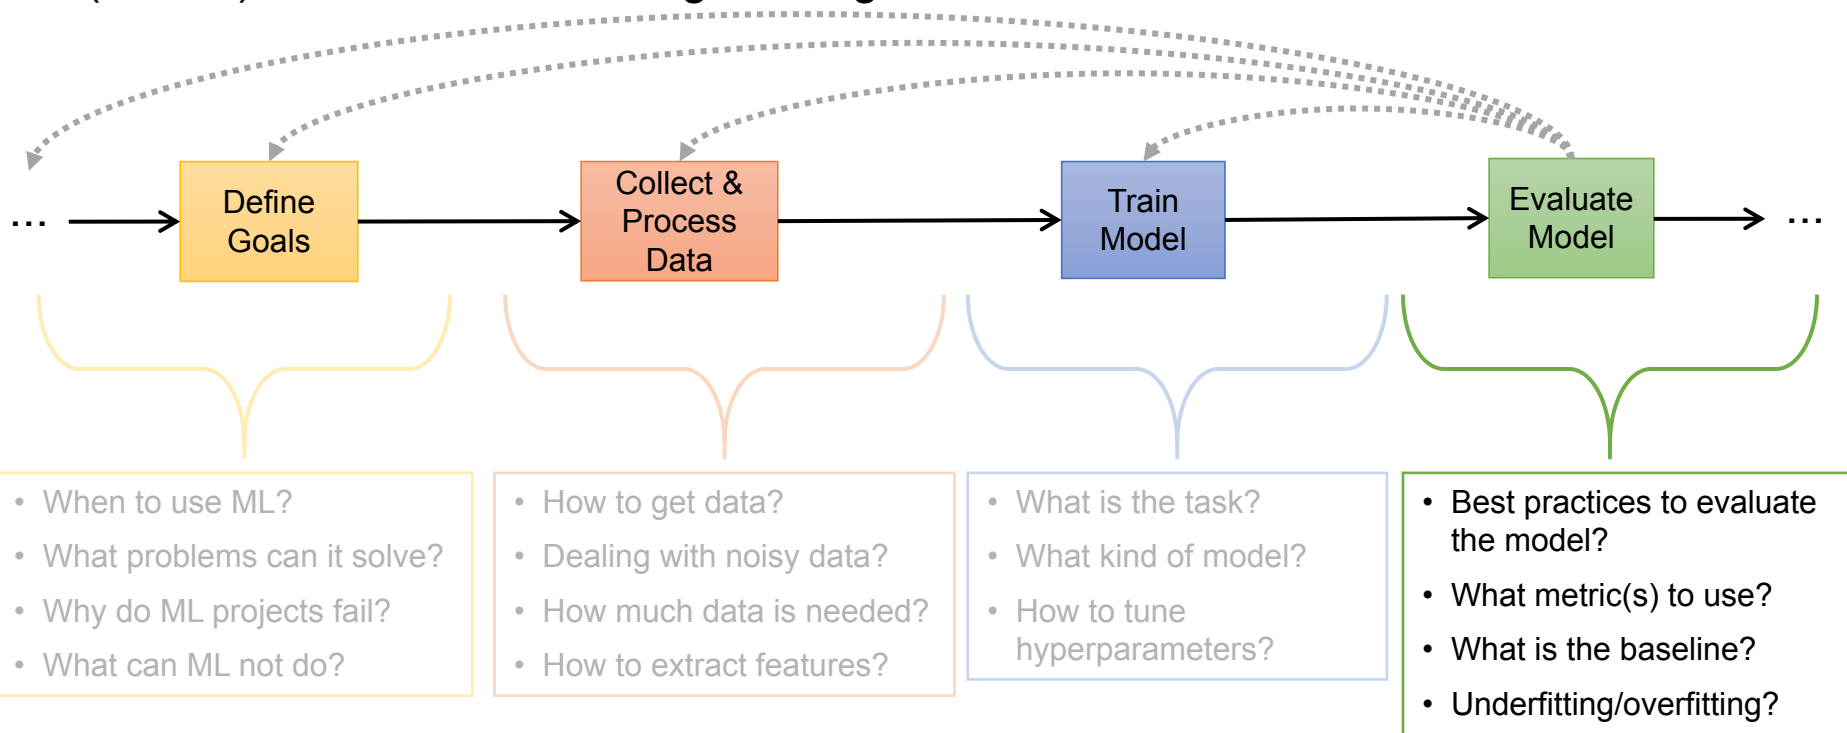
- ◆ These are not learned when we trained the model; the machine learning engineer sets those
 - ✿ E.g.: k in k NN is a hyperparameter; for SVM is the non-linearly separable case there is C (regularization hyperparameter)
- ◆ However, hyperparameters should be tuned
 - ✿ If we want the best model, we also need the best values of hyperparameters!

■ Hyperparameter tuning/optimization strategies

- ◆ **Grid search**: try all combinations of hyperparam values
 - ✿ For example: for hyperparam $a \in \mathcal{A}$ and $b \in \mathcal{B}$, try all pairs $(a, b) \in \mathcal{A} \times \mathcal{B}$
- ◆ **Random search**: given distribution of hyperparam values, we randomly sample from them
- ◆ Many others: e.g., bayesian hyperparameter optimization, evolutionary optimization, etc.

Machine Learning Engineering

■ (Partial) Workflow for ML Engineering:



Training, Test, Validation

- Before training a model (or looking at the data ideally)
 - ◆ Divide the dataset into **three disjoint parts**
 1. Training dataset
 2. Test dataset
 3. Validation dataset
- Why? Why do we need the validation dataset? Why not just training and test?
 - ◆ We need it for **hyperparameter optimization!**
 - ✿ Q: Why can't we use the training set for that?
 - ✿ Q: Why can't we use the test set for that?
- What proportion of the dataset to allocate to each?
 - ◆ (Rule of thumb:) For small datasets (i.e., < 100k examples): 70% training, 15% validation, 15% test
 - ◆ For large datasets (e.g., deep learning): 95% training, 2.5% validation, 2.5% test
 - ◆ For very small datasets (e.g., <1000 examples): check the raw numbers (e.g., how many examples is 10%?)
 - ◆ What if you don't have enough data to afford leaving some aside for validation/testing?
 - ✿ Use **k-fold validation**: divide the data into k equal parts, then train on $k-1$, test on the remaining part, repeat k times and average!

(Training) Data Leakage

■ Subtle failure more for ML: **data leakage**

- ◆ Occurs if the model is given access to information (at **training time**) that **would not be available** at **inference time**.
- ◆ **Row-wise** leakage (training example) or **column-wise** leakage
- ◆ Examples:
 - ✧ Duplicate data
 - ✧ Preprocessing leakage (e.g., premature feature engineering)
 - ✧ Improper hyperparameter tuning
 - ✧ Proxy attributes (e.g., want to predict `age` but `year_of_birth` is a feature)
 - ✧ Time leakage (e.g., time series data improperly split between train and test)
 - ✧ Etc.
- ◆ Major concern: **reproducibility** crisis
 - ✧ Ref: Kapoor and Narayanan. "*Leakage and the reproducibility crisis in machine-learning-based science*." Patterns 4, no. 9 (2023).

■ Bias

- ◆ Error due to **incorrect assumptions** in the model
- ◆ *Inability to capture the true relationship*
 - ✿ If a model is too simple to capture the true relationship between features and label/target, it will have high bias!
- ◆ **High bias** means **underfitting**!
- ◆ Terminology:
 - ✿ do not confuse this with the bias term in the parameters of a model (i.e., the intercept)

■ Variance

- ◆ Sensitivity to **small variation** in the training data
 - ✿ Think of training a model as a repeated randomized process
 - ✿ If the model is highly influenced by a few data points, then it has high variance (it models the random noise!)
- ◆ **High variance** means **overfitting**!

Bias and Variance

■ Bias


- ◆ Error due to **incorrect assumptions** in the model
- ◆ *Inability to capture the true relationship*

■ Variance


- ◆ Sensitivity to **small variations** in the training data

■ Ideally, we want: low bias **and** low variance


- ◆ Strategies to lower bias:
 - ✧ Increase model complexity
 - ✧ Use more features
- ◆ Strategies to lower variance:
 - ✧ Reduce model complexity
 - ✧ Use more training data




High bias



High variance



Low(er) bias &
low(er) variance



High bias &
High variance


Bias-Variance Tradeoff

- Generalization error (aka **out-of-sample error** or **risk**)

- ◆ Prediction error on *unseen* data
- ◆ Related to overfitting
 - ✿ If the model overfits, then the generalization error will be large

- Bias-Variance Tradeoff

- ◆ Generalization error: $\text{bias}^2 + \text{variance} + \text{irreducible error}$
 - ✿ For more details:
 - Geman et al. "Neural networks and the bias/variance dilemma." Neural computation (1992)
 - Kohavi et al. "Bias plus variance decomposition for zero-one loss functions." ICML, 1996.
- ◆ Why is it a tradeoff?
 - ✿ Increasing model complexity \Rightarrow lower bias
 - ✿ Decreasing model complexity \Rightarrow lower variance
 - ✿ Note: there has been some debate of whether this applies to neural networks
 - E.g.: see Neal et al. "A modern take on the bias-variance tradeoff in neural networks." arXiv, 2018.



- Most models can be regularized
 - ◆ Typically tuned through a **regularization constant** (hyperparameter)
 - ◆ Effect: lower variance at the cost of (slightly?) higher bias

- Regularization reduces model complexity
 - ◆ It decreases the **degrees of freedom** of the model
 - ✧ E.g.: for linear SVM, regularization controls the cost of misclassification in the loss function
 - ◆ Note: there are several types of regularization and regularization techniques

- If your model is overfitted
 - ◆ Regularization is (one of) the first things you should try

Measuring Bias & Variance

■ Key quantities:

- ◆ Error on training dataset
- ◆ Error on validation/test dataset
- ◆ To keep in mind: the **irreducible error**

■ Examples (classification):

◆ Assumptions:

- ✿ We measure the error using 1-accuracy
- ✿ Irreducible error is 0%

◆ Diagnoses:

1. Training error: 1%; validation error: 20% => low bias; high variance (**overfitted**)
2. Training error: 20%; validation error: 21% => high bias; low variance (**underfitted; generalizes well**)
3. Training error: 20%; validation error: 35% => high bias; high variance (**worst case**)
4. Training error: 1%; validation error: 2% => low bias; low variance (**best case**)

- In general, we do not know the **irreducible error**
 - ◆ Suppose the training error of a classifier is 20%
 - ◆ Q: Does the classifier have high bias (is it underfitted)?
 - ✧ **It depends what the irreducible error is!**

- It is critical to have an appropriate baseline!
 - ◆ Given a baseline, we can at least know if the model learned anything at all!
 - ◆ Baseline(s) for classification tasks
 - ✧ Random guessing:
 - If there c classes, the baseline accuracy is $1/c$ (baseline error is $1-1/c$)
 - ✧ Guessing the **mode** (most frequent class)
 - If q_i is the frequency of class i , then baseline error is $\min(1-q_i) = 1 - \max(q_i)$
 - ✧ If the problem is well-studied, use benchmarks as a baseline!
 - Note: *if humans can perform the task with almost 0% error, then the irreducible error is probably 0*

Next Time

- Friday (1/26): Exercise 2

- Upcoming:
 - ◆ **Homework 1** will be out today (due 2/2 by 11:59pm)