

# CAI 4104/6108 – Machine Learning Engineering: Ensembles


Prof. Vincent Bindschaedler

Spring 2024

# Reminder: Supervised Learning


## ■ Classification

- ◆ Task: predict the corresponding **label**
- ◆ Different types:
  - ✧ Binary classification: there are only two classes (0,1; +,-, etc.)
  - ✧ Multiclass: more than two classes
  - ✧ Multi-label: each instance can belong to more than one class
  - ✧ One-class: there is only one class, we want to distinguish it from everything else



## ■ Regression

- ◆ Task: predict the corresponding **value** (typically a real number) or **target**
  - ✧ E.g.: you want to predict a person's future income based on their high school GPA




## ■ Others:

- ◆ Sequence-to-sequence, similarity learning/metric learning, learning to rank, etc.


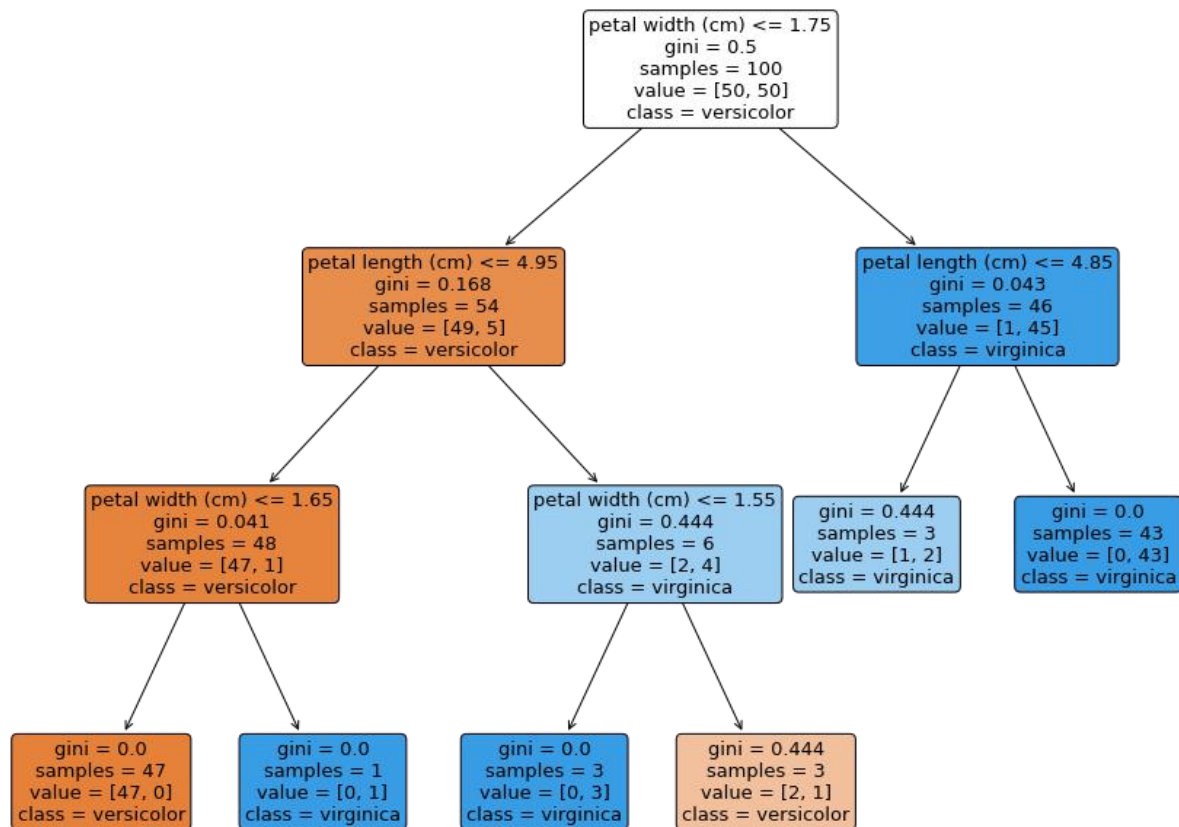
# Reminder: Decision Trees

## ■ A **decision tree** is

- ◆ An *acyclic* graph (i.e., a directed rooted tree) that can be used to make predictions
- ◆ **Nonparametric** model suitable for classification or regression
  - ✿ What is the other nonparametric model we have seen?
- ◆ Prediction:
  - ✿ Start at the root
  - ✿ Traverse the tree (branching according to feature values)
  - ✿ The leaf gives the predicted **label** or **value/target**
- ◆ How is the tree constructed from the training data?
  - ✿ There are many algorithms and many different kinds of decision trees!



# Reminder: Decision Tree Example



# Reminder: Overfitting & Regularization

- Decision trees make almost no assumption about the data
  - ◆ So unless we control complexity, the tree structure will be made to (over)fit the data!
  
- How do we avoid overfitting?
  - ◆ Prevent the tree from growing too deep (e.g., set a maximum depth)
  - ◆ Restrict splitting (e.g., set a minimum number of examples to split)
  - ◆ Pruning: after the tree is created, prune branches that do not significantly reduce the error
  
- Regularization hyperparameters
  - ◆ Example: Scikit-learn CART
    - ✧ `max_depth`: maximum depth of tree (default = unlimited)
    - ✧ `min_samples_split`: minimum number of examples to split a node (default = 2)
    - ✧ `min_samples_leaf`: minimum number of examples for a leaf (default = 1)

## ■ Motivating example:

- ◆ We have the following models (all for the same prediction task on the same data)
  - ✧ Model A: Linear SVM — 79% accuracy
  - ✧ Model B: SVM with RBF kernel — 81% accuracy
  - ✧ Model C: Nearest Neighbor (k=3) — 76% accuracy
  - ✧ Model D: Logistic Regression — 80% accuracy
  - ✧ Model E: Decision Tree — 78% accuracy
- ◆ Which model do we choose?
  - ✧ All of them? This is called an **ensemble** model.

## ■ Why would using multiple models be better than using just one?

## ■ How do we combine the models?

## ■ Voting Classifiers

- ◆ Given a feature vector, use every classifier to predict the label
- ◆ **Hard voting**: predict the majority label (*statistical mode* of all predictions)
- ◆ **Soft voting**: predict label with highest probability averaged over all classifiers
  - ✿ Only feasible for models that compute/estimate probabilities

## ■ Voting Regressors

- ◆ Predict the average (or median) of the prediction of all the regression models

- Q: Why would the average (or majority) of many models be better than a single model?
- Q: Is it better if the models are different or similar?
- Q: Do the models have to be any good?

# Ensemble Learning: Bagging

## ■ Bagging (**bootstrap aggregating**)

- ◆ Idea: train many different models of the same type
  - ✧ Each model is trained on a randomly sampled subset of the data
  - ✧ **Bagging**: sampling with replacement      **Pasting**: sampling without replacement
- ◆ Typically applied to decision trees,
  - ✧ But you could apply it to any type of classifier/regressor
- ◆ Aggregation function:
  - ✧ For classification: statistical mode
  - ✧ For regression: average
- ◆ Effect: lowers variance (and reduces overfitting) but bias is similar
- ◆ Variants
  - ✧ **Random subspaces**: pick random subsets of features (instead of examples)
  - ✧ **Random patches**: pick random subsets of features **and** of examples



## ■ Random Forests

- ◆ Bagging with Decision Trees
- ◆ But: when constructing a decision tree, use a random set of features to decide on best split!
  - ✿ This gives you more diversity among trees, which means an overall better model

## ■ Extremely Randomized Trees (Extra Trees)

- ◆ Like random forests but even more random! How?
  - ✿ When building the decision trees, pick a random threshold!

## ■ Note: there are other types of tree-based ensembles

- ◆ **Isolation Forests**: used mostly for anomaly detection
- ◆ **Embedding w/ Random Trees**: for unsupervised representation learning



# (Hypothesis) Boosting

## ■ Boosting

- ◆ Combine many **weak learners** into a strong learner
- ◆ Typically, we train the weak learners sequentially (each learn corrects errors of the previous one)
- ◆ Many variants but most popular are **Gradient Boosting** and **AdaBoost**

## ■ Gradient Boosting

- ◆  $h_0$ : base model (e.g., decision tree regressor)
- ◆  $h_1$ : model trained on **residual errors** of  $h_0$
- ◆  $h_2$ : model trained on **residual errors** of  $h_0 + h_1$
- ◆ Prediction:  $h_0(\mathbf{x}) + h_1(\mathbf{x}) + h_2(\mathbf{x})$

## ■ AdaBoost

- ◆  $h_0$ : base model (e.g., decision tree classifier)
- ◆  $h_i$ : model but with weights of examples misclassified by  $h_{i-1}$  increased

## ■ Stacking

- ◆ Alternative to voting classifiers / regressors
- ◆ Instead of averaging or majority voting, we train a model to do the aggregation
- ◆ **Meta model** (blender) is trained to combine the predictions
  
- ◆ Best practice:
  - ✧ Split the training data into two sets  $\mathcal{S}_1$  and  $\mathcal{S}_2$
  - ✧ Use  $\mathcal{S}_1$  to train the models
  - ✧ Make predictions on instances of  $\mathcal{S}_2$
  - ✧ Then use these predictions and the true labels/targets in  $\mathcal{S}_2$  to train the meta model

# Mixture of Experts (MoE)

- Recently reinvigorated idea in the deep learning era
  - ◆ MoE is now used often for state-of-the-art **Transformers** / **Large Language Models** (LLMs)
    - ✧ OpenAI's GPT-4 is believed to be a MoE
    - ✧ Mixtral 8x7B (from Mistral AI) is a MoE with 8 experts (7B parameters each)
- MoE Idea: divide-and-conquer strategy
  - ◆ Train a bunch of **experts**  $f_1, f_2, \dots, f_k$  (each **expert model** is usually a **neural network**)
  - ◆ **Weighting** (aka **gating**) function  $w_1(\mathbf{x}), w_2(\mathbf{x}), \dots, w_k(\mathbf{x})$ , that **depends on input  $\mathbf{x}$** 
    - ✧ The weight  $w_i(\mathbf{x})$  is the weight of  $f_i(\mathbf{x})$  — influence of expert  $i$  on the final prediction
    - ✧ Note: this is also learned from data (trained)
  - ◆ At **inference time** given input  $\mathbf{x}$ , the prediction is:  $\sum_i w_i(\mathbf{x}) f_i(\mathbf{x})$
- Many Variants:
  - ◆ Sparsely-gated MoE: use only the **top 1 (or 2)** experts according to the gating network

# Next Time

- Friday (2/9): Exercise
- Upcoming:
  - ◆ **Homework 2** will be out soon (due 2/14) by 11:59pm