

# CAI 4104/6108 – Machine Learning Engineering: Gradient Descent

Prof. Vincent Bindschaedler

Spring 2024

# Administrivia: Midterm

- **Reminder: Midterm** is coming up!
  - ◆ Monday **2/19** and Wednesday **2/21** during class time (10:40 - 11:30) in **FLG 0220**
    - ✿ Topics: everything until 2/16
  - ◆ Duration: 50 minutes
  - ◆ Schedule:
    - ✿ **CAI6108**: take the exam on **2/19**
    - ✿ **CAI4104**: take the exam on **2/21**
    - ✿ Lecture that week (topic: Unsupervised Learning) will be pre-recorded
  - ◆ Format: **closed-books**, (blank) scratch paper, and physical calculator are allowed (no phones!)
  - ◆ Questions: short answers + problems
  
- **Reminder: Sample Midterm** (Practice Questions):
  - ◆ On Canvas — 50 minutes (18 short answers) + 3 problems.
  - ◆ This is for you to practice. Do **not** overfit!

# Reminder: Metrics for Classification

## ■ Confusion Matrix

- ◆ Applicable to classification in general

## ■ Focus on the binary classification case:

- ◆ One of the classes is designated as “*positive*” (the other is “*negative*”)
- ◆ **False positives** are **Type I** errors
  - ✿ Think of them as “**false alarms**”
- ◆ **False negatives** are **Type II** errors
  - ✿ Think of them as “**missed detections**”
- ◆ **Prevalence** (aka **base rate**):
  - ✿ Proportion of positive examples

## ■ Baselines?

- ◆ Statistical mode prediction
- ◆ Random guessing

		Actual	
		+	-
Predicted	+	True Positive (TP)	False Positive (FP)
	-	False Negative (FN)	True Negative (TN)

- **Accuracy**:  $(TP + TN) / (TP + FP + TN + FN)$
- **Recall**:  $TP / (TP + FN)$ 
  - ◆ Also called: **True positive rate** (TPR), Sensitivity
- False negative rate (FNR):  $FN / (TP + FN) = 1 - TPR$
- **False positive rate** (FPR):  $FP / (FP + TN)$
- True negative rate (TNR):  $TN / (FP + TN) = 1 - FPR$ 
  - ◆ Also called Specificity and Selectivity
- **Precision**:  $TP / (TP + FP)$ 
  - ◆ Also called: **Positive predictive value** (PPV)
- False discovery rate (FDR):  $FP / (TP + FP) = 1 - Precision$

# Reminder: Precision-Recall Tradeoff

## ■ Decision Threshold / Function



- ◆ (In most cases) when you train a classification model, you actually train a **family of classifiers**!
  - ✿ The model assigns **scores** (or probabilities) to examples
  - ✿ Make predictions based on scores using a specific **decision threshold**

## ■ The Tradeoff

- ◆ The threshold provides a tradeoff between Type I and Type II errors
- ◆ A popular way to express this tradeoff is **Precision** versus **Recall**
  - ✿ We need to choose between high precision and high recall
  - ✿ Q: What if we need to achieve precision above a specific value (e.g., 90%)?

## ■ Optimizing and Statistficing

- ◆ One way of navigating the tradeoff is to set a cutoff for precision or recall
- ◆ E.g.: Pick the model with precision  $\geq 95\%$  that maximizes recall



# Reminder: Performance Curves

## ■ Receiver Operating Characteristics (ROC)

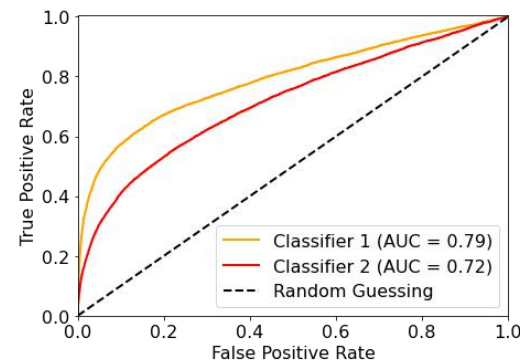

- ◆ Plots true positive rate (TPR) versus false positive rate (FPR)
  - ✿ How? Vary the threshold
- ◆ Each point on the curve (a valid pair (TPR,FPR)) is a valid tradeoff point
  - ✿ E.g.: **Equal Error Rate** (EER) — point where FPR and FNR are equal

## ■ Area Under Curve (AUC)

- ◆ This is exactly the area under the ROC curve
  - ✿ AUC = 0 means the worst possible classifier
  - ✿ AUC = 0.5 is a random classifier
  - ✿ AUC = 1.0 is a perfect classifier

## ■ Note: there are other performance curves

- ◆ For example: **Detection Error Tradeoff** (DET) curves
  - ✿ False positive rate (FPR) versus the false negative rate (FNR), usually a log-log plot



# Base Rate Fallacy

## ■ Example:

- ◆ Suppose you have a very accurate classifier (e.g., 99% accurate) to predict whether a person suffers from a specific disease D from features of their blood
- ◆ Q: Should we test everyone in the world? Why or why not?
  - ✿ It depends on the **base rate**!

## ■ Base rate fallacy / base rate neglect

- ◆ Error in reasoning: confusing a classifier's prior probability of correct prediction and the posterior probability of a true positive
- ◆ Suppose we have a classifier that has a false positive rate of 2% (i.e., 2% of the time it predicts '+' when the true label is '-') and a true positive rate of 100% (i.e., it never fails to detect '+' instances)
- ◆ What is the probability that if the classifier predicts '+' the true label is in fact '+'?
  - ✿ If the base rate is 0.5 (i.e., 50% of instances are '+') then it is: ~98%
  - ✿ If the base rate is 0.001 (i.e., 0.1% of instances are '+') then it is: **~4.8%**

# Base Rate Fallacy

## ■ Base rate fallacy / base rate neglect

- ◆ Suppose we have a classifier that has a false positive rate of 2% (i.e., 2% of the time it predicts '+' when the true label is '-') and a true positive rate of 100% (i.e., it never fails to detect '+' instances)
- ◆ What is the probability that if the classifier predicts '+' the true label is in fact '+'?
  - ✿ If the base rate is 0.5 (i.e., 50% of instances are '+') then it is: ~98%
  - ✿ If the base rate is 0.001 (i.e., 0.1% of instances are '+') then it is: ~4.8%

		Actual	
		+	-
Predicted	+	5000	100
	-	0	4900

$$\Pr(\text{Actual } + \mid \text{Predicted } +) = 5000/5100 \approx 0.98$$

		Actual	
		+	-
Predicted	+	10	200
	-	0	9790


$$\Pr(\text{Actual } + \mid \text{Predicted } +) = 10/210 \approx 0.0476$$

## ■ Note: there are other (similar) errors in reasoning

- ◆ Examples: Prosecutor's fallacy, Simpson's paradox, etc.

# Gradient Descent

- To solve an optimization problem, we can use **Gradient Descent** (or one of its variants)
  - ◆ Generic **iterative procedure** to update the parameters based on the **gradient**
- We want to minimize the loss function:  $L(\theta)$ 
  - ◆ Here  $\theta$  are the **parameters** (weights)
  - ◆  $\eta$  is the **learning rate** — size of steps we take towards the minimum
  - ◆  $\nabla_{\theta} L(\theta)$  is the **gradient of the loss** with respect to the parameters





# Gradient Descent

- Inputs: dataset  $\mathbf{X}$ ,  $\mathbf{y}$  ; model  $f$ , loss function  $L$  ; learning rate  $\eta > 0$  ; max number of iterations  $k$
- Output: parameter vector  $\theta$
- Procedure:

- ◆ Initialize  $\theta^{(1)}$  randomly

- ◆ For  $j = 1, 2, \dots, k-1$

- ✧ Take a batch of data  $\mathbf{X}_B, \mathbf{y}_B$


- ✧ Compute the gradient  $\mathbf{g} = \nabla_{\theta} L(f(\mathbf{X}_B; \theta^{(j)}), \mathbf{y}_B)$

- ✧ Update parameters:  $\theta^{(j+1)} = \theta^{(j)} - \eta \mathbf{g}$

- ◆ Output  $\theta^{(k)}$

This is the **gradient of the loss**  $L$  given current model parameters taken **with respect to the parameters**  $\theta$

Update **parameters** by taking steps of size  $\eta$  in the **opposite direction** of the **gradient**




- Q: Are we guaranteed to converge to optimal parameters  $\theta^*$ ?


# Reminder: Convex Sets & Convexity

## ■ Convex Sets:


Set  $\mathcal{X}$  is **convex** if for any  $a, b \in \mathcal{X}$  the line  $\lambda a + (1-\lambda)b \in \mathcal{X}$  for  $\lambda \in [0,1]$



convex





not convex




convex

## ■ Convexity:


Function  $f$  on a convex set  $\mathcal{X}$  is **convex** if for any  $x, y \in \mathcal{X}$  and  $\lambda \in [0,1]$ :  
$$\lambda f(x) + (1-\lambda)f(y) \geq f(\lambda x + (1-\lambda)y)$$



strictly convex



convex




not convex

# Types of Gradient Descent

## ■ Procedure:

- ◆ Initialize  $\theta^{(1)}$  randomly
- ◆ For  $j = 1, 2, \dots, k-1$ 
  - ✧ Take a batch of data  $X_B, y_B$
  - ✧ Compute the gradient  $\mathbf{g} = \nabla_{\theta} L(f(X_B; \theta^{(j)}), y_B)$
  - ✧ Update parameters:  $\theta^{(j+1)} = \theta^{(j)} - \eta \mathbf{g}$
- ◆ Output  $\theta^{(k)}$



## ■ Batch Gradient Descent (GD)

- ◆ In each iteration: calculate the gradient of the entire dataset [i.e.,  $(X_B, y_B) = (X, y)$ ]

## ■ Stochastic (“on-line”) Gradient Descent

- ◆ In each iteration, randomly pick a single example and calculate its gradient [i.e.,  $(X_B, y_B) = (x_i, y_i)$ ]

## ■ Mini-batch Stochastic Gradient Descent (“SGD”)

- ◆ Randomly shuffle the entire dataset and process it in mini-batches steps [i.e.,  $(X_B, y_B)$  is a mini-batch of  $m$  examples]
  - ✧ One epoch is the number of iterations to process the entire dataset
- ◆ What we typically use in practice: has stochasticity but does not require computing the gradient over the entire dataset

# Important Considerations

- Major consideration: **learning schedule**

- ◆ Way to set the **learning rate** in each iteration

- ✧ E.g.: constant, time-based decay, step decay, exponential decay, etc.


- Variants / **Optimizers**:

- ◆ AdaGrad, RMSProp, Adam, etc.

- **Feature scaling**: we **should** rescale features before doing gradient descent

- ◆ What happens if we don't?

- ✧ Convergence may be slow...



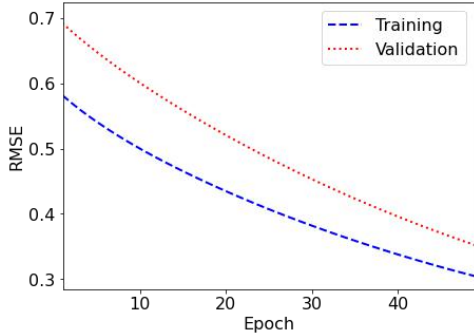

# Illustrated Example

- Let's train a linear regression model using gradient descent
  - ◆ Loss: Mean Squared Error (MSE)
  - ◆  $\text{MSE}(\mathbf{w}, b) := 1/n \sum_i [\mathbf{w} \mathbf{x}_i + b - y_i]^2 = 1/n \sum_i [\boldsymbol{\theta} \mathbf{x}_i - y_i]^2 = \text{MSE}(\boldsymbol{\theta})$ 
    - ✿ Here:  $\boldsymbol{\theta} = (b, \mathbf{w})$  and we extend  $\mathbf{X}$  to incorporate a constant feature ( $\mathbf{x}_{i,0}=1$  for all  $i$ )
- Gradient step:  $\boldsymbol{\theta}^{(j+1)} = \boldsymbol{\theta}^{(j)} - \eta \nabla_{\boldsymbol{\theta}} \text{MSE}(\boldsymbol{\theta})$ 
  - ◆  $\eta$  is the **learning rate** (hyperparameter)
- What is the gradient of  $\text{MSE}(\boldsymbol{\theta})$ ?
  - ◆  $\nabla_{\boldsymbol{\theta}} \text{MSE}(\boldsymbol{\theta}) = 2/n \mathbf{X}^T (\boldsymbol{\theta} \mathbf{X} - \mathbf{y})$
  - ◆ Q: What about the gradient for ridge regression?



# Learning Rate?

## ■ How do we set the learning rate?

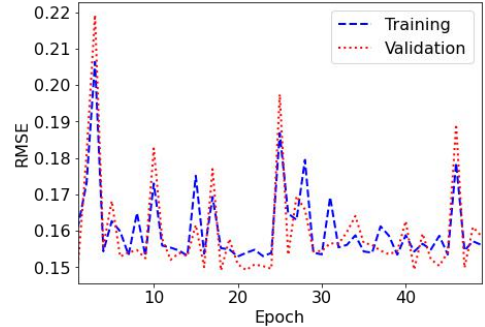

$\eta = 0.005$  (too small)



$\eta = 0.05$  (just right)



$\eta = 0.5$  (too large)



# Next Time

- Friday (2/16): Exercise and (or) Q&A
- Upcoming:
  - ◆ **Midterm** on 2/19 and 2/21