

CAI 4104/6108 – Machine Learning Engineering: Unsupervised Learning

Prof. Vincent Bindschaedler

Spring 2024


Unsupervised Learning

- Learning from **unlabeled data** (we must **discover patterns** in the data)

- Problems / Tasks

- ◆ Density Estimation

- ✳ **Model** the (unknown) **probability distribution function** (pdf) from which the data is drawn



- ◆ Outlier Detection

- ✳ **Identify** data points that **do not belong** in (i.e., are very different from) the (training) data

- ◆ Clustering

- ✳ **Group** similar data points together in “**clusters**”

- ◆ Dimensionality Reduction (& Visualization)

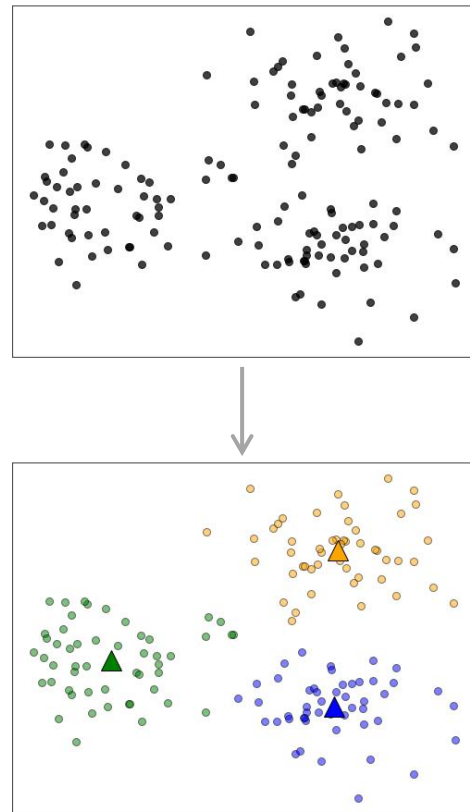
- ✳ **Transform** dataset into a **low(er)-dimensional** representation in a way that preserves useful information

- ◆ Association Rule Mining

- ✳ **Discover** interesting **relationships** between variables in the data

Clustering

- Goal: **assign** data points to clusters in the optimal way
 - ◆ What does optimal mean?
 - ◆ How many clusters?
- Dataset: Matrix \mathbf{X} ($n \times m$); **no labels**
- Clustering algorithms
 - ◆ k -means clustering
 - ◆ DBSCAN
 - ◆ Hierarchical clustering
- Q: Do we need a “holdout” set?
 - ◆ Yes! The clustering could **overfit**.
 - ◆ So we should train - test split or train - test - val split



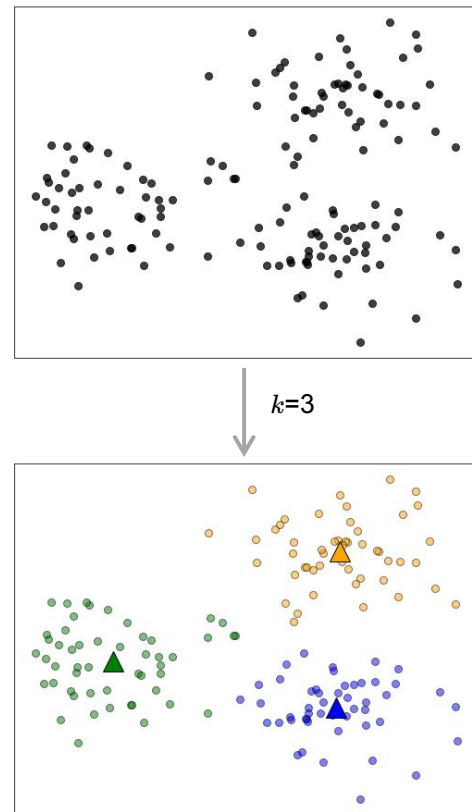
Clustering: K-Means

■ K-Means Clustering

- ◆ **Greedy iterative** algorithm to assign points to clusters
- ◆ Input: dataset \mathbf{X} , integer k (**hyperparameter** — number of clusters)
- ◆ Objective function: $L = \sum_i ||\mathbf{x}_i - c(\mathbf{x}_i)||^2$
 - ✧ Here $c(\mathbf{x}_j)$ is the centroid of the cluster that \mathbf{x}_j is assigned to.
- ◆ Algorithm:
 - ✧ Initialization: random centroid for each cluster
 - ✧ Do until no change in clustering:
 - Assignment: Assign each data point to the cluster with the closest centroid
 - Update: Recalculate centroids from points assigned to each cluster


■ Q: How do we determine k ?

- ◆ Guess or try a bunch of different options
- ◆ Heuristics: **Silhouette** method or **Elbow** method
- ◆ Tune it as you would any other hyperparameter (if you have an objective metric to evaluate the quality of a clustering — e.g. co-membership test)
- ◆ Or: Use a different algorithm that does not require specifying k (e.g., DBSCAN)



■ Curse of Dimensionality

- ◆ Working with high dimensional data is challenging
- ◆ For example: what do we do if each example in our dataset has millions of features?
 - ✿ Not only would training be slow, but model performance may suffer. Why?
- ◆ What can we do?
 - ✿ Increase the amount of training data to compensate for the large number of features
 - Not feasible in practice (in most cases)
 - ✿ If the dimensionality of the data is **artificial**, then we should be able to reduce dimension without losing information



■ Dimensionality Reduction Techniques

- ◆ Projections (e.g., PCA)
 - ✿ Project the data into a lower dimensional subspace by exploiting feature correlations
- ◆ Manifold Learning (e.g., LLE, Isomap)
 - ✿ Assume the data lies on some manifold. Model that manifold to reduce dimensionality
- ◆ Note: another motivation for reducing dimensionality is visualization

Principal Components Analysis (PCA)

■ Principal Components Analysis (PCA)


◆ Method for linear transformation onto a new system of coordinates

- ✿ Invented by Pearson in 1901
- ✿ The transformation is such that the **principal components** (coordinate vectors) capture the greatest amount of variance given the previous components

◆ Algorithm:

- ✿ Given data matrix \mathbf{X} ($n \times m$)
 - **Mean-center** it: subtract the mean of each feature
- ✿ Compute covariance matrix $\mathbf{X}^T\mathbf{X}$ ($m \times m$)
- ✿ The eigendecomposition of gives principal components
 - We can use Singular Value Decomposition (SVD)
 - Matrix \mathbf{W} of eigenvectors is the transformation matrix (the i^{th} column is the i^{th} principal component)
 - Eigenvalue λ_i gives the variance of i^{th} principal component


◆ Note: we can do PCA on the correlation matrix instead of covariance matrix



Principal Components Analysis (PCA)

■ PCA Transformation / Dimensionality Reduction

- ◆ Given data matrix \mathbf{X} ($n \times m$), transformation matrix \mathbf{W} ($m \times m$), and eigenvalues λ
 - ✱ $\mathbf{Z} = \mathbf{X} \mathbf{W}$ is the transformed data onto the PCA coordinate system
- ◆ If we want to reduce dimensionality, we can keep only the **first k principal components**
 - ✱ Think of k as a hyperparameter
 - ✱ Let \mathbf{W}_k is the transformation matrix with only the first k columns ($m \times k$)
 - ✱ Transformed data: $\mathbf{Z}_k = \mathbf{X} \mathbf{W}_k$ [here is a $n \times k$ matrix]
- ◆ How much of the variance is explained?
 - ✱ Variance explained: $(\lambda_1 + \lambda_2 + \dots + \lambda_k) / \sum_i \lambda_i$
- ◆ How do we determine k ?
 - ✱ Same way we tune other hyperparameters
 - ✱ Or: set k to explain a fixed portion of the variance (e.g., 95%)



Principal Components Analysis (PCA)


■ PCA Transformation / Dimensionality Reduction

◆ **Inverse transformation:** can we transform the data back to the original space?

✧ Yes: $\mathbf{X}' = \mathbf{Z}_k \mathbf{W}_k^\top$ where \mathbf{W}_k^\top is a $k \times m$ matrix

✧ Here: \mathbf{X}' is a $(n \times m)$ matrix (note: some information was lost in the transformation unless $k = m$)

◆ What does this actually look like?





Kernel PCA

- PCA is a linear decomposition/transformation

- ◆ What if we need *non-linear* dimensionality reduction?



- ✿ We can use the **kernel trick** (the same one we used for SVM); this is called **Kernel PCA**

Swiss Roll



Linear kernel (= PCA)

RBF kernel ($\gamma=0.05$)




Polynomial kernel
(degree 3)

■ Manifold Learning

- ◆ Given a dataset $X = \{x_1, x_2, \dots, x_n\}$ embedded in \mathbb{R}^m (each point $x_i \in \mathbb{R}^m$)
- ◆ Find a mapping of the dataset X into a dataset $Z = \{z_1, z_2, \dots, z_n\}$ embedded in \mathbb{R}^p ($z_i \in \mathbb{R}^p$) for some integer $p < m$
- ◆ Such that (informally) Z preserves the local geometry of X
 - ✧ For example if x_i and x_j are close (according to some distance metric), then z_i and z_j are also close
 - ✧ So really we want to preserve the neighborhood structure

■ Multidimensional Scaling (MDS)

- ◆ Compute the Euclidean distance d_{ij} between any two points x_i and x_j
- ◆ We want to minimize: $\sum_{i < j} [\|z_i - z_j\| - d_{ij}]^2$
 - ✧ The minimization here is over z_1, z_2, \dots, z_n
- ◆ What if we don't care about (approximately) preserving large distances?
 - ✧ Then we could ask to only preserve distances only if two points are neighbors




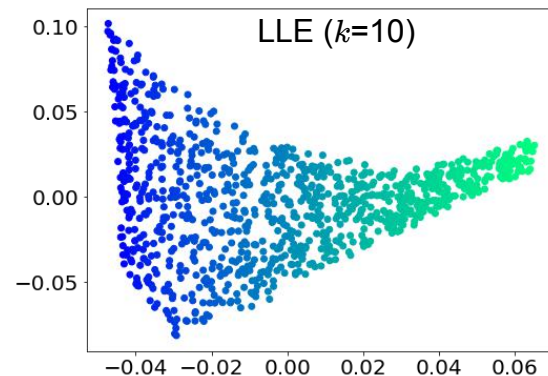
Manifold Learning

■ Algorithms

- ◆ Different algorithms aim to preserve different “local” properties
- ◆ Multidimensional Scaling (MDS)
 - ✧ Preserve distances between points
- ◆ **Locally Linear Embedding (LLE)**
 - ✧ Express each data point as a linear combination of its closest neighbors
- ◆ **Isomap**
 - ✧ Form a graph where data points are connected to their closest neighbors
 - ✧ Aims to preserve geodesic distance (i.e., shortest path distance)
- ◆ **t-distributed Stochastic Neighbor Embedding (t-SNE)**
 - ✧ Tries to keep similar data points close together, dissimilar data points far apart
 - ✧ Mostly used for visualization

■ Q: Does feature scaling matter?

- ◆ Yes, we should rescale the data (to ensure all features have the scale)
- ◆ Why? Methods are based on nearest-neighbors search



Next Time

- Friday (2/23): Exercise
- Upcoming:
 - ◆ Homework 2 due 2/23