

CAI 4104/6108 – Machine Learning Engineering: Neural Networks

Prof. Vincent Bindschaedler

Spring 2024

History of Neural Networks


■ (Artificial) **Neural Networks**

- ◆ Large class of models / architecture
- ◆ **Loosely** inspired by the biology of our brains

■ Timeline

- ◆ 1943: McCulloch and Pitts. "A logical calculus of the ideas immanent in nervous activity." Bulletin of Mathematical Biophysics.
- ◆ 1958: Perceptron algorithm
- ◆ 1960s: backpropagation derived by many researchers independently
- ◆ 1980s: application of backpropagation to multi-layer neural networks by Rumelhart, Hinton and Williams (1986) and Yann Lecun in his PhD thesis (1987)
- ◆ 2010s: deep learning revolution
 - ✿ greater availability of data; more computational power; techniques to overcome difficulty of training deep neural networks; fast implementations on GPUs, etc.

Neural Network Terminology




A Simple Neural Network

■ Consider a single neuron / unit

- ◆ The model is $h_{w,b}(x) = f(w \cdot x + b)$
 - ✿ What if we take f to be the identity function?
 - That is: $f(z) = z$
 - ✿ What if we take f to be the **sigmoid** / **logistic** function?
 - That is: $f(z) = 1/(1+e^{-z})$

■ The **Perceptron**

- ◆ Invented by Frank Rosenblatt in 1957
 - ✿ "The Perceptron—a perceiving and recognizing automaton". Report 85-460-1. Cornell Aeronautical Laboratory
- ◆ A different neuronal architecture called a threshold linear unit (TLU)
 - ✿ No bias term
 - ✿ With a **step** activation function. For example:
 - $\text{heaviside}(z) = 0$ if $z \leq 0$, 1 otherwise ($z \geq 1$) ; or $\text{sign}(z)$



Components

■ Types of Layers

- ◆ Dense (i.e., fully-connected)
- ◆ Convolutional
- ◆ Recurrent

■ Activation Functions


- ◆ **Identity / Linear** (or none): $f(z) = z$
- ◆ **Sigmoid**: $f(z) = 1/(1+e^{-z})$
- ◆ **TanH**: $f(z) = (e^z - e^{-z}) / (e^z + e^{-z})$
- ◆ **ReLU**: $f(z) = \max(0, z)$
- ◆ **Softmax**: $f(z_j) = \exp(z_j / T) / \sum_i \exp(z_i / T)$

✿ Note: in that case the activation function is over an entire layer, not a single unit

■ Loss

- ◆ Whatever you like (e.g., squared error loss) as long as it's differentiable

✿ Note: make sure the loss function and activation function of the output layer are consistent with each other!



Examples & Special Cases

■ Single neuron?

◆ Linear regression

- ✿ One layer neural network with a single neuron with a linear activation function and MSE as loss function

◆ (Binary) Logistic regression

- ✿ One layer neural network with a single neuron with a sigmoid activation function and binary cross-entropy as loss function

■ Multi-Layer Perceptron (MLP)

◆ Input layer (passthrough)

◆ One or more hidden layers

◆ Output layer

■ Neural network with multiple layers all with linear activation

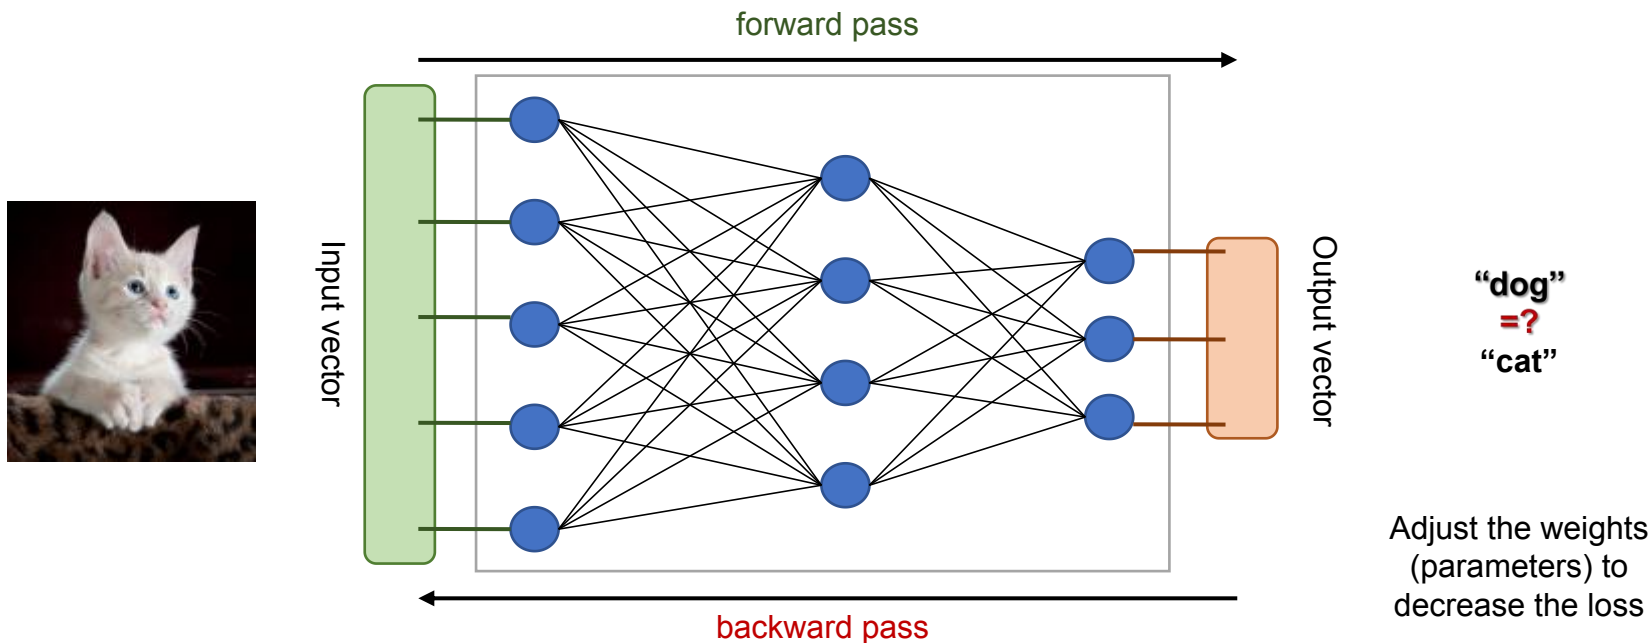
◆ Q: Good idea?

- ✿ **No. A linear function of a linear function is still a linear function!**

Training Neural Networks

■ Backpropagation:

- ◆ Reverse pass to measure error and propagate error gradient backwards in the network
 - ✿ Using gradient descent to update the parameters



Backpropagation

■ Seminal Paper:


- ◆ “*Learning representations by back-propagating errors.*”
Rumelhart, Hinton, and Williams. Nature 1986.

■ Terminology


- ◆ **Backpropagation**: how to compute the gradients efficiently
- ◆ **Gradient descent**: how to update the parameters to minimize the loss given the gradient

■ Algorithm: given a mini-batch B

- ◆ Compute the **forward pass** for the mini-batch B saving the intermediate results at each layer
- ◆ Compute the **loss** on the mini-batch B (compares output of network to labels/targets → **error**)
- ◆ **Backwards pass**: computes the per-weight gradients (**error contribution**) layer by layer
 - ✱ This is done using the **chain rule** (if z depends on y and y depends on x : $dz/dx = dz/dy \cdot dy/dx$)
- ◆ (Stochastic) gradient descent: update the weights based on the gradients



Backpropagation: Illustration



output: $a^{(l)}_j = f^{(l)}(z^{(l)}_j)$

$$z^{(l)}_j = \sum_i w^{(l)}_{ij} a^{(l-1)}_i + b^{(l)}_j$$

- How should we **update** $w^{(l)}_{ij}$?
 - ◆ Based on $\partial L / \partial w^{(l)}_{ij}$ where L is the loss
 - ◆ How to compute $\partial L / \partial w^{(l)}_{ij}$?
 - ✳ Activation output $f^{(l)}$ is a function of $z^{(l)}_j$ and $z^{(l)}_j$ is a function of $w^{(l)}_{ij}$
 - ✳ Chain rule: $\partial L / \partial w^{(l)}_{ij} = \partial L / \partial z^{(l)}_j \cdot \partial z^{(l)}_j / \partial w^{(l)}_{ij}$
 - ✳ $\partial z^{(l)}_j / \partial w^{(l)}_{ij} = a^{(l-1)}_i$
 - ◆ If we let $\delta^{(l)}_j = \partial L / \partial z^{(l)}_j$, we see that $\delta^{(l-1)}_i$ depends on $\delta^{(l)}_j$
 - ✳ So we can compute the gradients **right to left** (i.e., **backwards**)

When To Use Neural Networks?

- You have a complex learning task or complex data
 - ◆ For many problems, neural networks provide the best performance (for some tasks performance is better than humans)
 - ✧ E.g.: image classification / captioning tasks
 - ✧ E.g.: speech recognition
 - ✧ E.g.: natural language modeling

- When should you NOT use a neural network?
 - ◆ If you can solve your problem with a simple model (e.g., linear model, decision tree, SVM)
 - ◆ If you do not have **a lot** of data
 - ✧ Neural networks (esp. deep neural networks) require a lot of data to achieve good performance
 - ◆ You need an explainable/interpretable model
 - ✧ Note: there exists techniques to explain decisions from neural networks

Universality of Neural Network

■ Universal Approximation Theorems



- ◆ *(Feed-forward) Neural networks can approximately represent any function*

- ◆ **Arbitrary width; bounded depth:**

 - ✧ True even if we have a single hidden layer as long as it can have arbitrarily many units

- ◆ **Bounded width; arbitrary depth:**

 - ✧ True even if we have layers of bounded width, as long as the network can have arbitrarily many layer



Neural Network Architecture?


■ Pitfall: inconsistent activation function of output layer with the loss function

◆ Examples:

- ✿ Multiclass classification with cross-entropy loss, softmax activation for output layer => **Okay**
- ✿ Regression with MSE as loss, tanh as activation for output layer => **Fail**
- ✿ Regression with MSE as loss, linear activation for output layer => **Okay**

■ Tip: the “funnel”

- ◆ For supervised learning we typically have large input feature vectors and small output vectors
 - ✿ We should make the network look like a funnel



■ Example: Multiclass classification with 10 classes and m=100 input features.

◆ The network could look like this:

- ✿ (Input, hidden layer 1, hidden layer 2, hidden layer 3, output layer)
 - 100, 64, 32, 16, 10
- ✿ Activations:
 - Output: Softmax
 - Elsewhere: ReLU

Next Time

- Wednesday (2/28): Lecture
- Upcoming:
 - ◆ Homework 3 will be out soon