


# CAI 4104/6108 – Machine Learning Engineering: Training Neural Networks

Prof. Vincent Bindschaedler

Spring 2024

- **Midterm** (update on grading)
  - ◆ We should be done grading later this week
  - ◆ Grades will most likely be out (sometime) next week
  
- **No class** on Friday (3/1)
  - ◆ Exercise 7 will be **pre-recorded** (available on Canvas)

# Reminder: Neural Network Terminology



# Reminder: A Simple Neural Network

## ■ Consider a single neuron / unit

◆ The model is  $h_{w,b}(x) = f(w \cdot x + b)$

- ✿ What if we take  $f$  to be the identity function?
  - That is:  $f(z) = z$
- ✿ What if we take  $f$  to be the **sigmoid** / **logistic** function?
  - That is:  $f(z) = 1/(1+e^{-z})$


## ■ The **Perceptron**

◆ Invented by Frank Rosenblatt in 1957

- ✿ "The Perceptron—a perceiving and recognizing automaton".  
Report 85-460-1. Cornell Aeronautical Laboratory

◆ A different neuronal architecture called a threshold linear unit (TLU)

- ✿ No bias term
- ✿ With a **step** activation function. For example:
  - $\text{heaviside}(z) = 0$  if  $z \leq 0$ ,  $1$  otherwise ( $z \geq 1$ ) ; or  $\text{sign}(z)$



# Reminder: Components

## ■ Types of Layers

- ◆ Dense (i.e., fully-connected)
- ◆ Convolutional
- ◆ Recurrent

## ■ Activation Functions


- ◆ **Identity / Linear** (or none):  $f(z) = z$
- ◆ **Sigmoid**:  $f(z) = 1/(1+e^{-z})$
- ◆ **Tanh**:  $f(z) = (e^z - e^{-z}) / (e^z + e^{-z})$
- ◆ **ReLU**:  $f(z) = \max(0, z)$
- ◆ **Softmax**:  $f(z_j) = \exp(z_j / T) / \sum_i \exp(z_i / T)$

✿ Note: in that case the activation function is over an entire layer, not a single unit

## ■ Loss

- ◆ Whatever you like (e.g., squared error loss) as long as it's differentiable

✿ Note: make sure the loss function and activation function of the output layer are consistent with each other!



# Reminder: Backpropagation

## ■ Seminal Paper:


- ◆ “*Learning representations by back-propagating errors.*”  
Rumelhart, Hinton, and Williams. Nature 1986.

## ■ Terminology

- ◆ **Backpropagation**: how to compute the gradients efficiently
- ◆ **Gradient descent**: how to update the parameters to minimize the loss given the gradient

## ■ Algorithm: given a mini-batch B



- ◆ Compute the **forward pass** for the mini-batch B saving the intermediate results at each layer
- ◆ Compute the **loss** on the mini-batch B (compares output of network to labels/targets → **error**)
- ◆ **Backwards pass**: computes the per-weight gradients (**error contribution**) layer by layer
  - ✱ This is done using the **chain rule** (if  $z$  depends on  $y$  and  $y$  depends on  $x$ :  $dz/dx = dz/dy \cdot dy/dx$ )
- ◆ (Stochastic) gradient descent: update the weights based on the gradients



# Reminder: Universality of Neural Network

## ■ Universal Approximation Theorems

- ◆ *(Feed-forward) Neural networks can approximately represent any function*
- ◆ **Arbitrary width; bounded depth:**
  - ✧ True even if we have a single hidden layer as long as it can have arbitrarily many units
- ◆ **Bounded width; arbitrary depth:**
  - ✧ True even if we have layers of bounded width, as long as the network can have arbitrarily many layer



# Neural Network Architecture?


## ■ Pitfall: inconsistent activation function of output layer with the loss function

### ◆ Examples:

- ✿ Multiclass classification with cross-entropy loss, softmax activation for output layer => **Okay**
- ✿ Regression with MSE as loss, tanh as activation for output layer => **Fail**
- ✿ Regression with MSE as loss, linear activation for output layer => **Okay**

## ■ Tip: the “funnel”

- ◆ For supervised learning we typically have large input feature vectors and small output vectors
  - ✿ We should make the network look like a funnel



## ■ Example: Multiclass classification with 10 classes and m=100 input features.

### ◆ The network could look like this:


- ✿ (Input, hidden layer 1, hidden layer 2, hidden layer 3, output layer)
  - 100, 64, 32, 16, 10
- ✿ Activations:
  - Output: Softmax
  - Elsewhere: ReLU



# Deep Neural Networks

## ■ What is a **deep neural network**?

- ◆ Any neural network with two or more hidden layers
- ◆ Nowadays, the best neural networks architectures for many applications & problems are deep
  - ✿ E.g.: AlexNet (2012) has 8 layers, ResNet18 has 18 layers, GPT-2 has 48 layers




## ■ Challenges:

- ◆ Endless options for the network **architecture/topology**
  - ✧ E.g.: number of layers; units per layer; connections between units; activation functions; weight initialization method
- ◆ Hyperparameters related to learning:
  - ✧ E.g.: optimizer, learning rate, decay/momentum, (mini)batch size, number of epochs, etc.

## ■ Rules of Thumb:

- ◆ Number of hidden layers
  - ✧ **Deep > shallow**: For the same number of parameters, more hidden layers is better than wider layers
  - ✧ Why? **Parameter efficiency**
- ◆ Number of units in each layer
  - ✧ **Funnel** approach: make the network look like a funnel
  - ✧ **“Stretch pants”** approach: make hidden layers wider than what you need and then regularize (e.g., dropout)
    - Ref: Vanhoucke’ Udacity course on Deep Learning




# Architecture & Hyperparameters Tuning

## ■ Challenges:

- ◆ Endless options for the network **architecture/topology**
  - ✿ E.g.: number of layers; units per layer; connections between units; activation functions; weight initialization method
- ◆ Hyperparameters related to learning:
  - ✿ E.g.: optimizer, learning rate, decay/momentum, (mini)batch size, number of epochs, etc.

## ■ Rules of Thumb:

- ◆ Activation functions
  - ✿ Hidden layers => ReLU (or ReLU variants)
    - Faster to compute than alternatives; Gradient descent less likely to get “stuck”
  - ✿ Output layer:
    - **Multiclass classification**: *softmax*
    - **Binary classification** or **multilabel**: *sigmoid*
    - **Regression**: *linear* (no activation function)



## ■ Challenges:

- ◆ Endless options for the network **architecture/topology**
  - ✿ E.g.: number of layers; units per layer; connections between units; activation functions; weight initialization method
- ◆ Hyperparameters related to learning:
  - ✿ E.g.: optimizer, learning rate, decay/momentum, (mini)batch size, number of epochs, etc.


## ■ Rules of Thumb:

- ◆ Learning rate:
  - ✿ Start with a low value (e.g., 0.000001) then multiply by 10 each time and train for a few epochs; once training diverges you have gone too far
- ◆ Optimizer: use **RMSProp** or **Adam**
- ◆ Batch size:
  - ✿ **Small batch** approach: e.g., 32, 64, 100
  - ✿ **Large batch** approach: the largest size that fits in your GPU's RAM (e.g., 8192) **and** use learning rate warmup
- ◆ Number of epochs/iterations: use **early stopping**



# Vanishing Gradient & Exploding Gradient

## ■ Vanishing/Exploding Problems:

- ◆ Gradient vector becomes very small (**vanishing gradient**) or very large (**exploding gradient**) during **backpropagation**
  - ✿ Difficult to update weights of lower/earlier layers => Training does not converge
- ◆ Instance of a more general problem: **unstable gradients**
  - ✿ Layers (of a deep neural network) learn at very different rates



Source: <https://cv-tricks.com/keras/understand-implement-resnets/>





# Vanishing Gradient & Exploding Gradient

## ■ Vanishing/Exploding Problems:

- ◆ Gradient vector becomes very small (**vanishing gradient**) or very large (**exploding gradient**) during **backpropagation**
  - ✿ Difficult to update weights of lower/earlier layers => Training does not converge
- ◆ Instance of a more general problem: **unstable gradients**
  - ✿ Layers (of a deep neural network) learn at very different rates

## ■ Mitigations:

- ◆ **Weight initialization method**
- ◆ **Non-saturating activation functions**
- ◆ **Batch normalization**
- ◆ **Gradient clipping** (for exploding gradient)
- ◆ **Skip-connections** (CNNs)
- ◆ ...



# Vanishing Gradient & Exploding Gradient

## ■ Weight Initialization Methods:

### ◆ Some ways to initialize the weights can make gradients unstable

- ✿ This held back efforts to train deep neural networks in the 2000s

### ◆ **Glorot initialization** (aka Xavier initialization):

- ✿ Seminal paper: Xavier Glorot and Yoshua Bengio. "*Understanding the difficulty of training deep feedforward neural networks.*" In AISTATS, 2010.
- ✿ Let  $n_{\text{in}}$ : number of inputs,  $n_{\text{out}}$ : number of outputs  $n_{\text{avg}} = (n_{\text{in}} + n_{\text{out}})/2$
- ✿ Gaussian (for sigmoid activation): mean 0, variance  $\sigma^2 = 1/n_{\text{avg}}$
- ✿ Uniform (for sigmoid activation): in  $[-r, r]$  where  $r = [3/n_{\text{avg}}]^{1/2}$  [default for Keras]

### ◆ **He initialization**:

- ✿ Kaiming He et al. "*Delving deep into rectifiers: Surpassing human-level performance on imagenet classification.*" In ICCV, 2015.
  - ✿ Gaussian (for ReLU): mean 0, variance  $\sigma^2 = 2/n_{\text{avg}}$
- ◆ Note: biases are initialized to 0

# Vanishing Gradient & Exploding Gradient

## ■ Non-saturating activation functions

- ◆ Activation functions like ReLU behave better than sigmoid and tanh

- ◆ But ReLU has one problem: **dying ReLUs**

- ✧ A neuron can “die” when the weighted sums of its input are negative (for all examples in the training data)

- ✧ In this case, ReLU would always output 0



- ◆ ReLU variants:

- ✧ LeakyReLU<sub>a</sub>(z) = max{az, z} (e.g., a = 0.01)

- ✧ ELU<sub>a</sub>(z) = z if z ≥ 0 and a(e<sup>z</sup> - 1) otherwise (z < 0) (e.g., a = 1)

- ✧ There is also SELU (Scaled ELU)

- ✧ These will not let neurons die because they can output negative values



## ■ So what to use in practice?

- ◆ My advice: use ReLU in most cases; if you have extra time use ELU (network will be slower)

- ◆ Although: ELU > leaky ReLU > ReLU > Tanh > Sigmoid



# Next Time

- Friday (3/1): **No Class**
  - ◆ *[Pre-recorded]* Exercise 7
- Upcoming:
  - ◆ **Homework 3** will be out soon