

CAI 4104/6108 – Machine Learning Engineering: Convolutional Neural Networks

Prof. Vincent Bindschaedler

Spring 2024

Administrivia: Project

- **Project Proposals** are due **3/27** on Canvas (by 11:59pm)
 - ◆ Create a group (4 or 5 students)
 - ✿ Canvas: Go to **People** → **Project Groups**
 - ◆ Identify a **task**, **dataset**, and some **baselines** (& metrics)
 - ◆ Fill in the project proposal template and submit the PDF on Canvas

- **Task & Dataset**: your choice
 - ◆ There **cannot** be overlap with other groups
 - ◆ Baselines and metrics have to make sense
 - ◆ **No** trivial projects please!
 - ✿ No toys datasets, MNIST classification has been done to death, etc.

Administrivia: Project


■ Proposal

- ◆ Length: 1 - 2 pages (use template)
- ◆ Structure:
 - ✧ Follow instructions in template
- ◆ **Writing counts! (try to be clear and concise)**

■ Submission on Canvas

- ◆ **Use provided Latex template!**
- ◆ **One per group**; write **all** group members names
- ◆ **Only one person** (point of contact) submits on Canvas
 - ✧ **Use the group functionality of Canvas**
 - ✧ You and **all** of your team mates must be in the **same project group** on Canvas

Reminder: Neural Network Terminology



Reminder: A Simple Neural Network

■ Consider a single neuron / unit

◆ The model is $h_{w,b}(x) = f(w \cdot x + b)$

- ✿ What if we take f to be the identity function?
 - That is: $f(z) = z$
- ✿ What if we take f to be the **sigmoid** / **logistic** function?
 - That is: $f(z) = 1/(1+e^{-z})$


■ The **Perceptron**

◆ Invented by Frank Rosenblatt in 1957

- ✿ "The Perceptron—a perceiving and recognizing automaton".
Report 85-460-1. Cornell Aeronautical Laboratory

◆ A different neuronal architecture called a threshold linear unit (TLU)


- ✿ No bias term
- ✿ With a **step** activation function. For example:
 - $\text{heaviside}(z) = 0$ if $z \leq 0$, 1 otherwise ($z \geq 1$) ; or $\text{sign}(z)$



Reminder: Deep Neural Networks

■ What is a **deep neural network**?


- ◆ Any neural network with two or more hidden layers
- ◆ Nowadays, the best neural networks architectures for many applications & problems are deep
 - ✧ E.g.: AlexNet (2012) has 8 layers, ResNet18 has 18 layers, GPT-2 has 48 layers



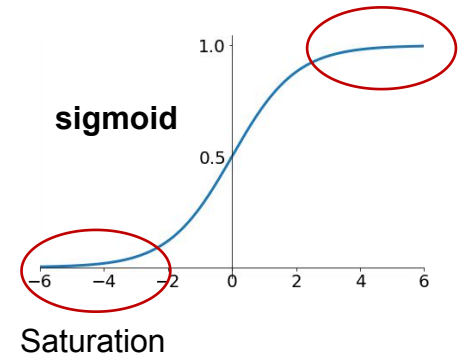

Reminder: Unstable Gradients

■ Vanishing/Exploding Problems:

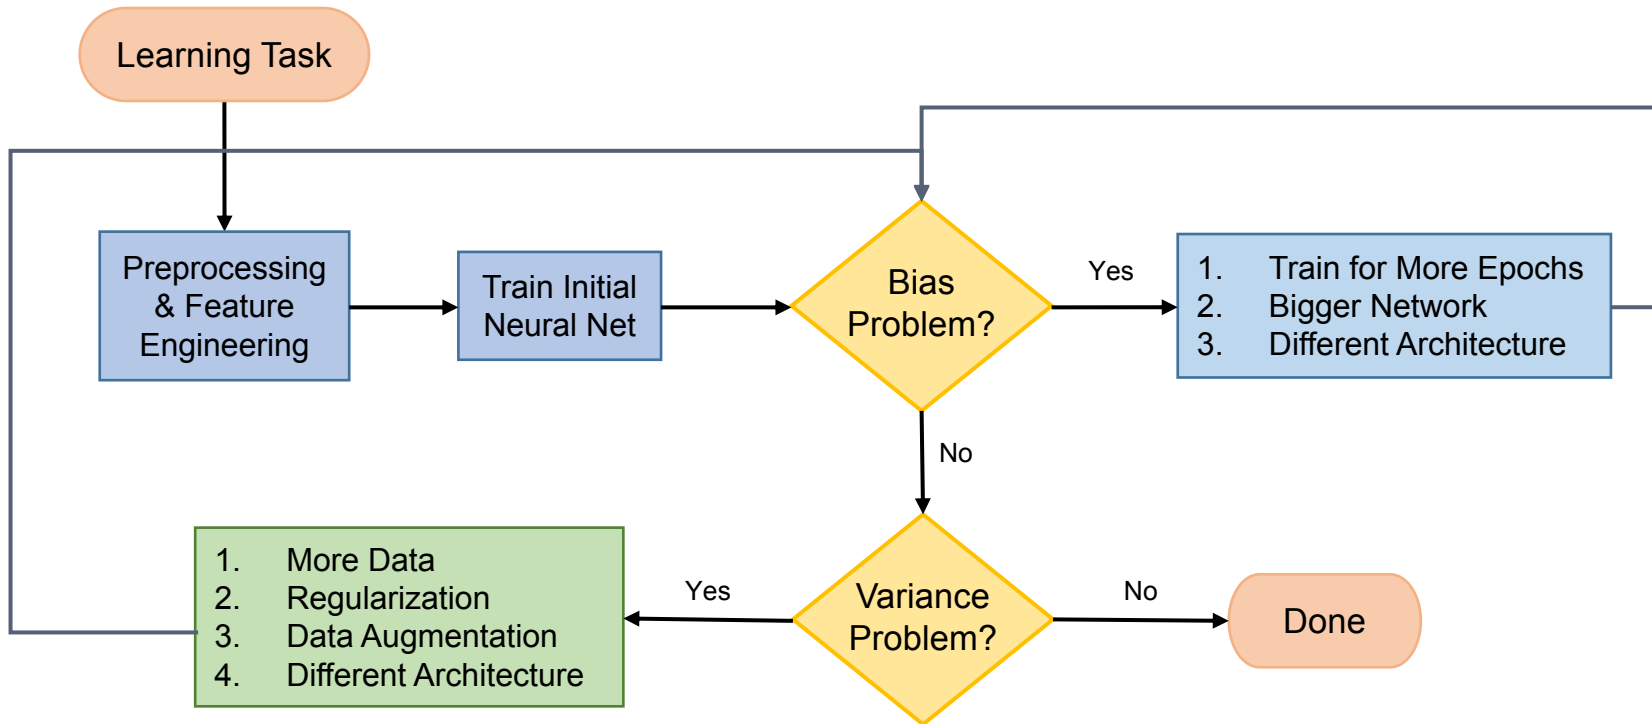
- ◆ Gradient vector becomes very small (**vanishing gradient**) or very large (**exploding gradient**) during **backpropagation**
 - ✿ Difficult to update weights of lower/earlier layers => Training does not converge
- ◆ Instance of a more general problem: **unstable gradients**
 - ✿ Layers (of a deep neural network) learn at very different rates



Source: <https://cv-tricks.com/keras/understand-implement-resnets/>



Reminder: Problem Solving



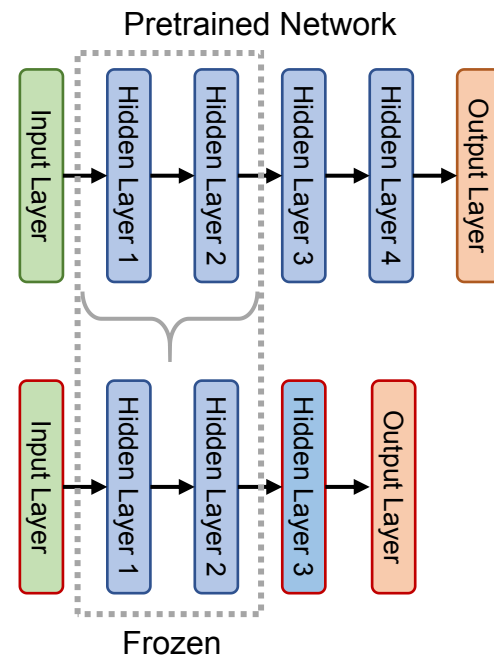
Reminder: Transfer Learning

■ Should you train a deep neural network from scratch?

- ◆ Not always. When possible you should use transfer learning:
 - ✧ Pick a **pre-trained deep neural network** in the same or related domain
 - ✧ Then **fine-tune** on the task you care about

■ Reusing an existing deep neural network

1. Pick some layers to reuse (typically the earlier layers)
2. **Freeze** these layers
 - ✧ This will set the corresponding parameters as **non-trainable**
 - Optimization: you can actually **cache** the outputs of frozen layers for every input
3. Add your own layers hidden layer(s)
4. Replace or discard upper layers
 - ✧ You should always discard the existing output layer and use your own



■ History:


- ◆ 1958: Hubel and Wiesel experiments on cats
 - ✿ Won the Nobel Prize in Physiology or Medicine (1981)
 - ✿ Insight: neurons in visual cortex have a **small local receptive field**
- ◆ 1998: LeCun et al. propose the LeNet-5 architecture

■ Convolutional Neural Networks:

- ◆ Architecture for neural networks using **convolutional layers**
 - ✿ **Convolutional layers**: each neuron/unit is only connected to a small number of neurons/units in the previous layer
 - ✿ Fewer neurons/units than fully-connected layers
- ◆ Well-suited to computer vision tasks or tasks on **image data**
 - ✿ Can also be applied to other tasks: for example some tasks in natural language processing
- ◆ Preeminent neural network architectures for many state-of-the-art applications
 - ✿ E.g.: self-driving cars, video classification, image search systems, etc.
- ◆ Remark: CNNs have high memory usage *during training*


Convolutional Neural Network Architecture

■ Example & Terminology:




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✿ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function



Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function




■ Remarks:


- ◆ The weights matrix remains the same throughout the convolution
 - ✧ There are only $f_w f_h + 1$ parameters for the filter (and it does not depend on the size of the input)
- ◆ Typically we have multiple filters per layer, so we get one feature map as output for each filter
- ◆ Output size of feature map depends on the size of the filter, stride, and padding strategy

Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function




(2,2) filter, **stride**=1,
padding "valid" (no padding)




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function




(2,2) filter, **stride=2**,
padding "valid" (no padding)




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function





(3,3) filter, **stride**=2,
padding "valid" (no padding)



Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✿ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function



Next Time

- Wednesday (3/20): Lecture

- Upcoming:
 - ◆ Homework 3 is due 3/20
 - ◆ Project