

CAI 4104/6108 – Machine Learning Engineering: Convolutional Neural Networks (2)

Prof. Vincent Bindschaedler

Spring 2024

■ History:

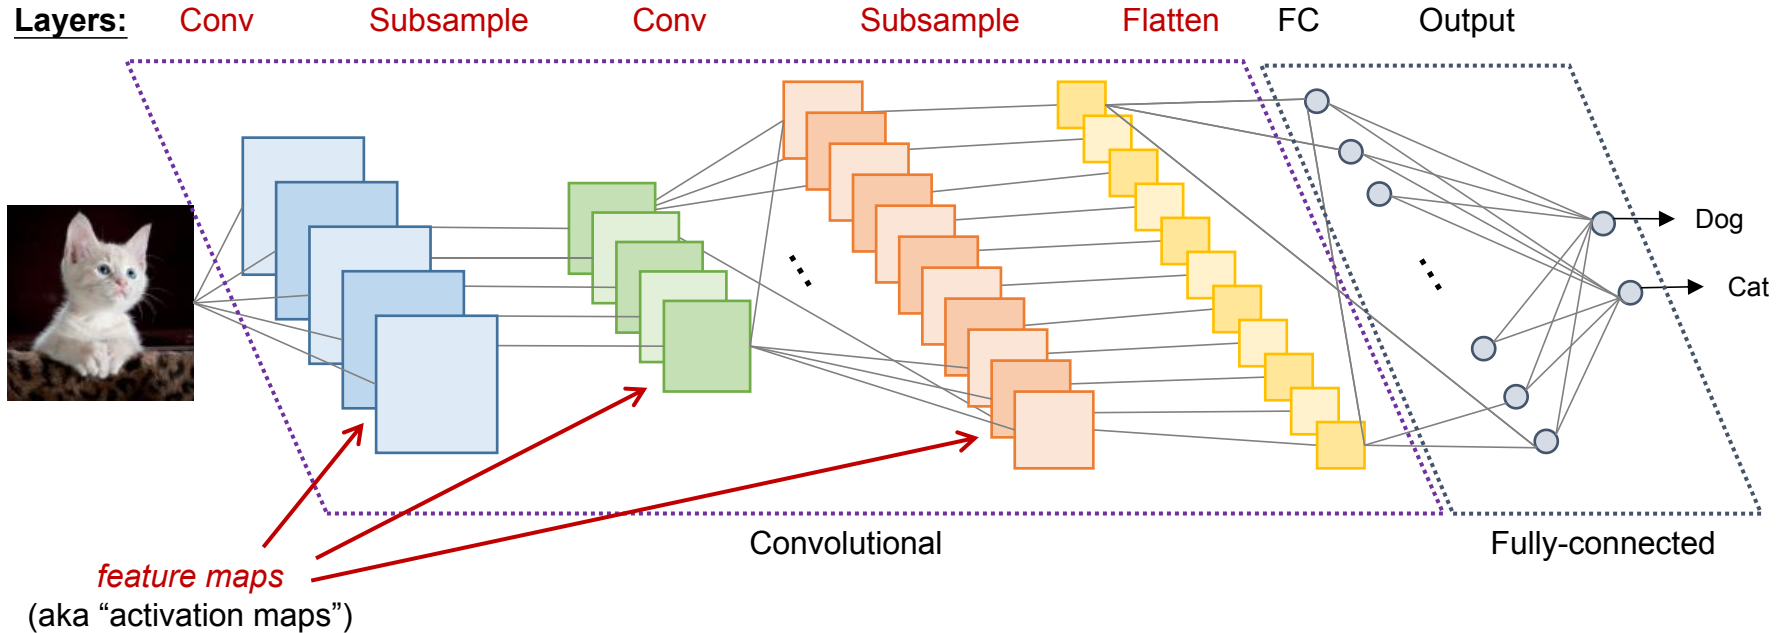
- ◆ 1958: Hubel and Wiesel experiments on cats
 - ✿ Won the Nobel Prize in Physiology or Medicine (1981)
 - ✿ Insight: neurons in visual cortex have a **small local receptive field**
- ◆ 1998: LeCun et al. propose the LeNet-5 architecture

■ Convolutional Neural Networks:

- ◆ Architecture for neural networks using **convolutional layers**
 - ✿ **Convolutional layers**: each neuron/unit is only connected to a small number of neurons/units in the previous layer
 - ✿ Fewer neurons/units than fully-connected layers
- ◆ Well-suited to computer vision tasks or tasks on **image data**
 - ✿ Can also be applied to other tasks: for example some tasks in natural language processing
- ◆ Preeminent neural network architectures for many state-of-the-art applications
 - ✿ E.g.: self-driving cars, video classification, image search systems, etc.
- ◆ Remark: CNNs have high memory usage *during training*


Reminder: CNN Architecture

■ Example & Terminology:




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✿ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function



Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function




■ Remarks:


- ◆ The weights matrix remains the same throughout the convolution
 - ✧ There are only $f_w f_h + 1$ parameters for the filter (and it does not depend on the size of the input)
- ◆ Typically we have multiple filters per layer, so we get one feature map as output for each filter
- ◆ Output size of feature map depends on the size of the filter, stride, and padding strategy

Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function




(2,2) filter, **stride**=1,
padding "valid" (no padding)




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function




(2,2) filter, **stride=2**,
padding "valid" (no padding)




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✿ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function





(3,3) filter, **stride**=2,
padding "valid" (no padding)




Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✧ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function



Convolutional Layer

- A convolutional layer has a set of **filters** (aka kernels)
 - ◆ Each filter **slides** (i.e., **convolves**) across the image (or previous layer's output) producing a **feature map**
 - ◆ The filter is represented by a $f_w \times f_h$ matrix of weights F and a bias b ; there is also an activation function
 - ✿ Applying the filter produces a **single output value** (real number) for each **sliding window**
 - ◆ Parameters: weight matrix F and bias b
 - ◆ Hyperparameters: **filter/kernel size** (f_w, f_h), **stride**, **padding strategy** ('valid' or 'same'), and **activation** function



How to calculate output (i.e., feature map) size?

$$\text{output_size} = \text{floor}[(V - K + 2P) / S] + 1$$


- V : input volume (e.g., width or height)
- K : kernel size
- P : padding (note: 0 for 'valid')
- S : stride

Subsampling / Pooling Layers

■ Subsampling Layers


- ◆ After one (or more) convolutional layers, we can have **subsampling** (aka “**pooling**”) layers to reduce the dimensions of feature maps
 - ✿ **Max pooling** layer: take the *maximum* value of the sliding window
 - ✿ **Average/mean pooling** layer: take the *mean* value of the sliding window
- ◆ Hyperparameter: **pooling size** (width, height) — for example: (2, 2) or (3, 3)
- ◆ Note: pooling layers do **not** have **any** parameters

Example: **Max pooling** (2, 2)



Example Architecture


■ Example & Terminology:



- How do we know what is a good CNN architecture for a problem?
 - ◆ There is no one-size-fits-all solution
 - ◆ Look at successful CNN architectures (e.g., LeNet-5, AlexNet, ResNet, etc) and adapt them to your problem

- Rules of thumb:
 - ◆ Avoid large filter sizes; stick to (2,2), (3,3) etc. But: the first layer can be larger (e.g., (5,5))
 - ◆ Use repetition / variants of the following patterns:
 - ✧ Pattern1: Conv, MaxPool
 - ✧ Pattern2: Conv, Conv, MaxPool
 - ◆ Use ReLU as the activation function (for convolutional layers)
 - ◆ The deeper you are in the network the more filters you want
 - ✧ For example: you could use 32 filters for the first conv layer, then 64 for the second, 128 for the third, etc.
 - ◆ Use **dropout** on the FC (dense) layer after you flatten

Example: LeNet-5



■ Source:


- ◆ LeCun et al. "Gradient-based learning applied to document recognition." Proceedings of the IEEE 86.11 (1998): 2278-2324.

■ Notes:

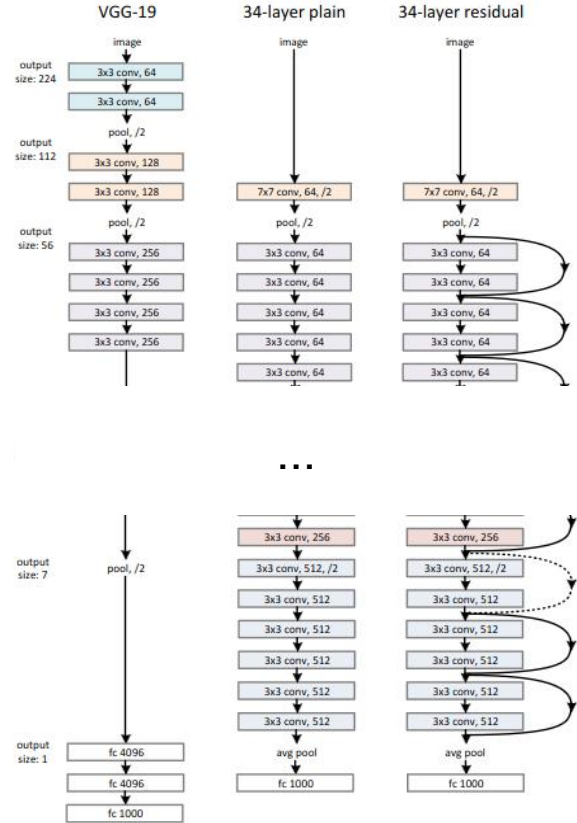
- ◆ All layers but the last use **tanh** as activation; nowadays we would use ReLU
- ◆ The subsampling layers are doing **average pooling**; nowadays we would use max pooling
- ◆ The output uses **RBF** activation; nowadays we would use softmax with crossentropy loss

Example: ResNets

Residual Learning building block



Source: He et al. "Deep residual learning for image recognition."
In Proceedings of the IEEE conference on computer vision and
pattern recognition, pp. 770-778. 2016.



Next Time

- Friday (3/22): Exercise

- Upcoming:
 - ◆ Homework 3 is due 3/20 (today)
 - ◆ Project