

CAI 4104/6108 – Machine Learning Engineering: Auto-Encoders & GANs


Prof. Vincent Bindschaedler

Spring 2024

Reminder: Deep Neural Networks

■ What is a **deep neural network**?

- ◆ Any neural network with two or more hidden layers
- ◆ Nowadays, the best neural networks architectures for many applications & problems are deep
 - ✿ E.g.: AlexNet (2012) has 8 layers, ResNet18 has 18 layers, GPT-2 has 48 layers



■ AutoEncoders

- ◆ Architecture combining an **encoder** network and a **decoder** network
- ◆ Learn efficient representations of the data
 - ✿ Each data point can be represented in the **latent space**
- ◆ Applications: dimensionality reduction and feature learning

■ Generative Adversarial Networks (GANs)

- ◆ Novel idea: **adversarial learning/training**


■ Generative models

- ◆ Some models can actually generate **new** data instances
- ◆ E.g.: some autoencoders, GANs, etc.

AutoEncoders

■ Encoder-Decoder network

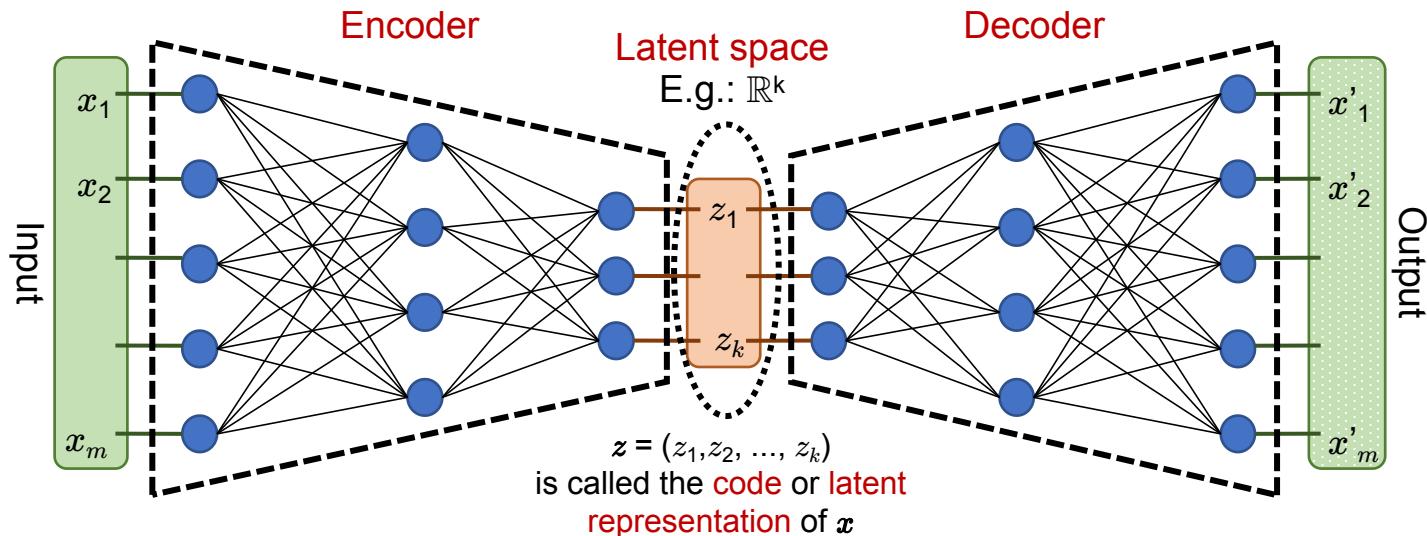
- ◆ Goal: learn to **reproduce** the input as output
- ◆ Constraints:
 - ✿ The **latent representation** (aka **codings**) is constrained (e.g., must have lower dimensionality than input)
- ◆ Effect: network must **learn** an **efficient** way to represent the information



AutoEncoders

■ Encoder-Decoder network


- ◆ Goal: learn to **reproduce** the input as output
- ◆ Constraints:
 - ✿ The **latent representation** (aka **codings**) is constrained (e.g., must have lower dimensionality than input)
- ◆ Effect: network must **learn** an **efficient** way to represent the information



Training AutoEncoders

■ Encoder-Decoder network

- ◆ Goal: learn to **reproduce** the input as output
- ◆ Constraints:
 - ✿ The **latent representation** (aka **codings**) is constrained (e.g., must have lower dimensionality than input)
- ◆ Effect: network must **learn** an **efficient** way to represent the information



Types of AutoEncoders

- **Deep** (aka **Stacked**) AutoEncoders:

- ◆ Multiple hidden layers for encoder and decoder
- ◆ Note: layers could be fully-connected, convolutional, recurrent, etc.

- **Sparse** AutoEncoders:

- ◆ Use a large bottleneck layer, but with a sparsity constraint (e.g., enforced through regularization)

- **Denoising** AutoEncoders:

- ◆ Add noise (typically Gaussian) to the input (or use dropout) to force the network to learn “robust” features and how to remove noise in the output

- **Variational** AutoEncoders:

- ◆ Probabilistic AutoEncoder, which makes it a generative model
- ◆ Idea: a data point is **encoded** as a mean μ and standard deviation σ
 - ✿ Then, we sample from a Gaussian with mean μ and standard deviation σ
 - ✿ Training loss: reconstruction loss (as before) + **KL-divergence** of latent space distribution and isotropic gaussian

- Many others...

Variational AutoEncoders

■ Seminal Paper


- ◆ Kingma and Welling. “Auto-Encoding Variational Bayes.” stat, 10. 2014

■ Probabilistic encoder/decoder

- ◆ Encoder maps an input x to a **distribution** in the latent space
 - ✧ Posterior $p(z|x)$
 - ✧ Approximate posterior $q(z|x)$
- ◆ Prior $p(z)$ over the latent space
 - ✧ Usually we choose Gaussian $\mathcal{N}(\mu, \sigma^2)$
- ◆ Likelihood $p(x|z)$

■ This is a **generative** model


- ◆ Q: How do we sample?



Training Variational AutoEncoders

■ Seminal Paper

- ◆ Kingma and Welling. “*Auto-Encoding Variational Bayes*.” stat, 10. 2014



■ How to train the model?

- ◆ Loss function: **reconstruction** loss (e.g., MSE or cross entropy) + **Kullback-Leibler divergence** between $p(z|x)$ and $q(z|x)$
 - ✧ Evidence Lower Bound (ELBO)
- ◆ How can we do backpropagation? The latent representation is **random**!?
 - ✧ **Reparameterization trick**: $z = \mu(x) + \epsilon\sigma(x)$
 - Here $\epsilon \sim \mathcal{N}(0, I)$ is an **external input**
 - Sometimes called “**stochastic backpropagation**”

- (Some) generative models allow us to:
 - ◆ Sampling — $x \sim p(x)$
 - ◆ We want to be able to **sample new instances** from the **learned** distribution
 - ◆ Density estimation — $p(x)$
 - ✧ We want to **estimate** $p(x)$ or compare $p(x_1)$ and $p(x_2)$
 - ◆ Learn representations — $z = \text{repr}(x)$
 - ✧ The representation can be used in **downstream tasks** (e.g., classification or regression)
 - ✧ And (in many cases) **reduce dimensionality** (e.g., use an AutoEncoder instead of PCA)

How Good Are AutoEncoders?

■ Variational Auto-Encoder (VAE)

Input



Reconstruction




Samples



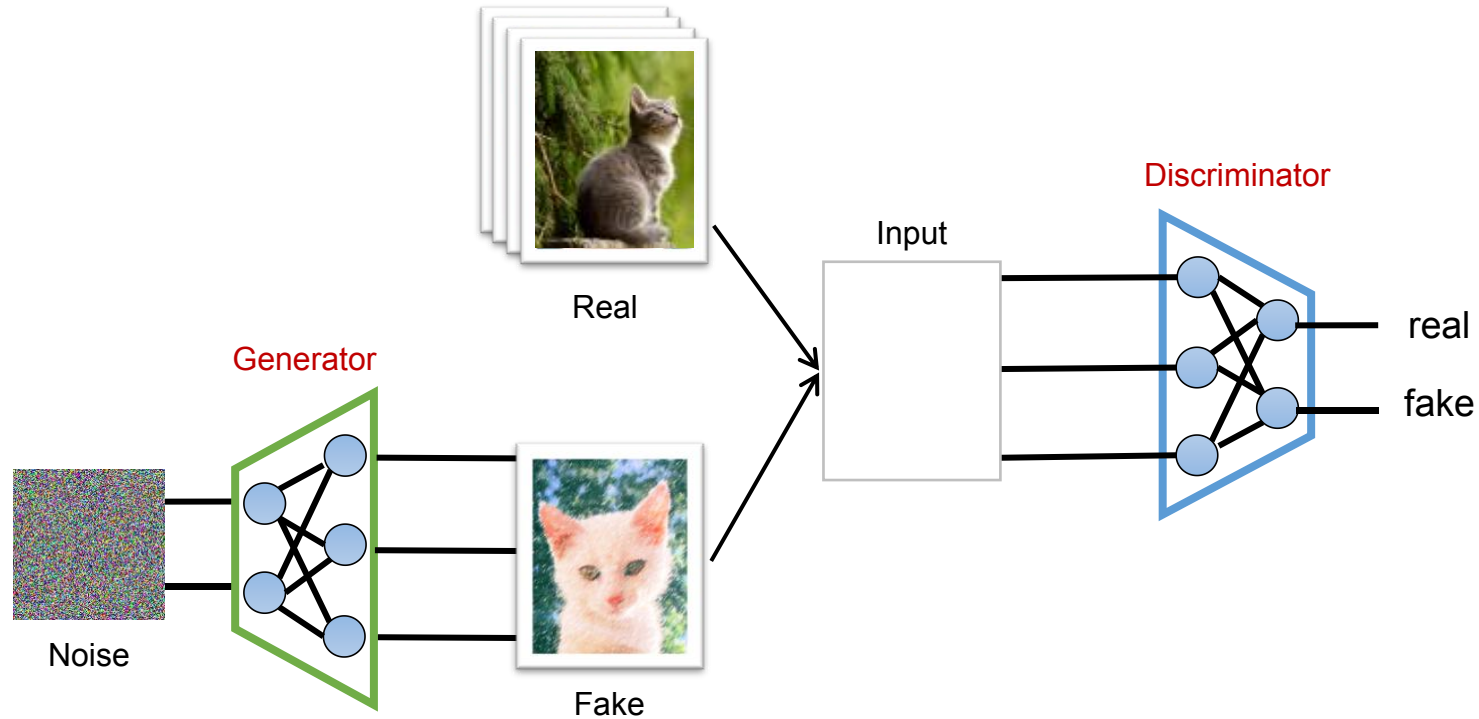
Source: Tolstikhin, Ilya, Olivier Bousquet, Sylvain Gelly, and Bernhard Schoelkopf. "Wasserstein auto-encoders." 2017.

Manipulating Latent Features

Source: Radford, Alec, Luke Metz, and Soumith Chintala.
"Unsupervised representation learning with deep convolutional
generative adversarial networks." arXiv, 2015.



Generative Adversarial Networks (GANs)



Generative Adversarial Networks

■ Origins:

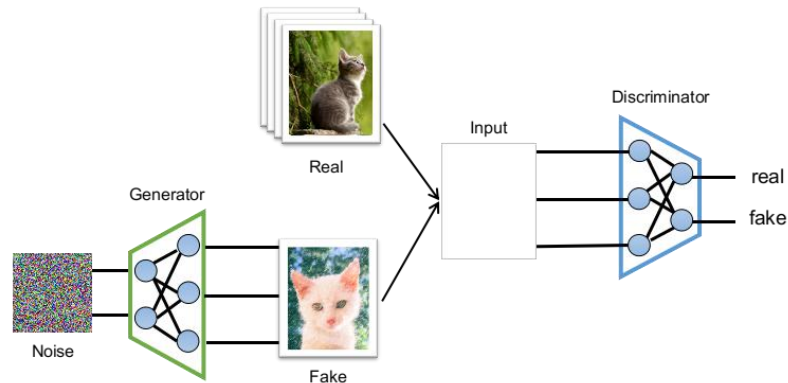
- ◆ Goodfellow et al. "*Generative Adversarial Nets.*" in NeurIPS, 2014.

■ Generator:

- ◆ Takes **random noise** from some distribution (e.g., gaussian) and produces a data point
- ◆ Trained using "feedback" from the discriminator

■ Discriminator:

- ◆ Given a data point predict **real** (1) or **fake** (0)
 - ✧ **Real**: data points taken from the dataset
 - ✧ **Fake**: data points produced by the generator



Training GANs

■ Challenges:

- ◆ GANs are notoriously difficult to train
- ◆ Generator and discriminator need to learn together at roughly the same pace
 - ✱ Otherwise, the training process will fail

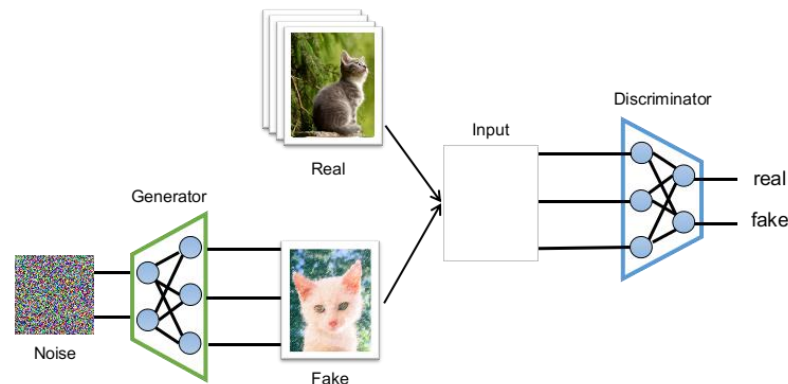
■ (Informal) training loop (for each epoch):

◆ Discriminator:

- ✱ Take k real data points (label 1)
- ✱ Run the generator to produce k fake data points (label 0)
- ✱ Train the discriminator on those $2k$ data points

◆ Generator:

- ✱ Freeze the weights of the discriminator (why?)
- ✱ Run the generator to produce k fake data points
- ✱ Give them to the discriminator ***pretending they are real***
- ✱ Backpropagate and update the weights!



How Good Are These Models?

■ DCGAN




*Source: Radford, Alec, Luke Metz, and Soumith Chintala.
"Unsupervised representation learning with deep convolutional
generative adversarial networks." arXiv, 2015.*



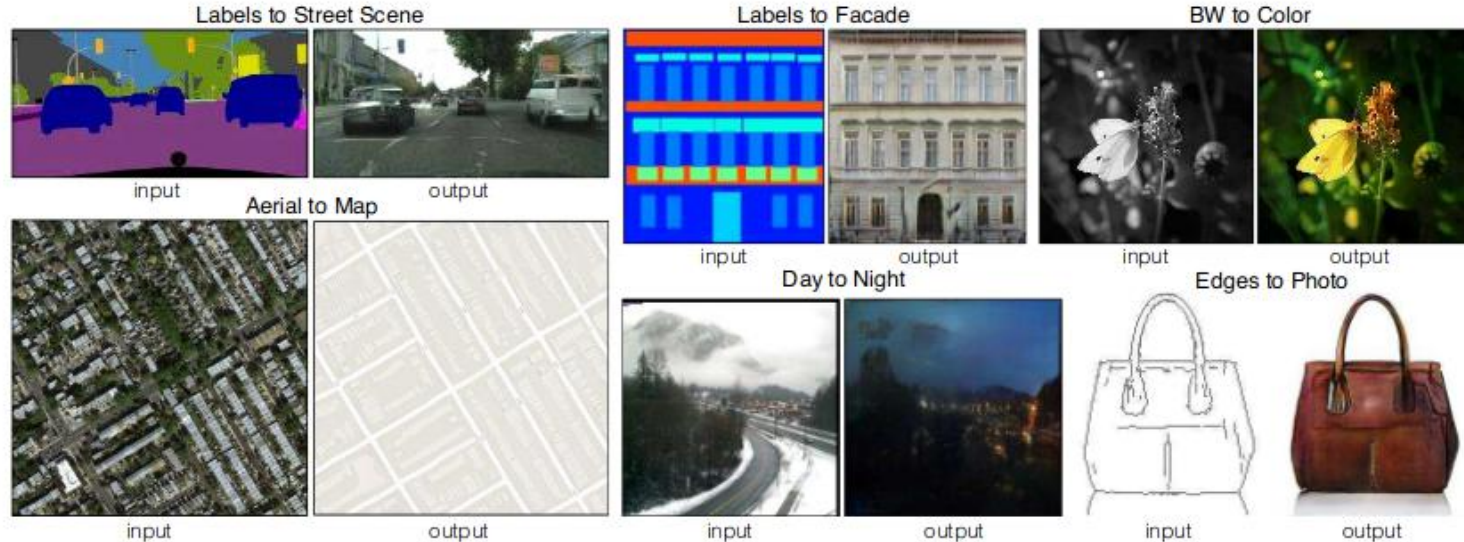
GANs are Improving

Source: Karras, Tero, Timo Aila, Samuli Laine, and Jaakko Lehtinen. "Progressive growing of gans for improved quality, stability, and variation." ICLR, 2018.



Other Applications of Generative Models

■ Image-to-Image translation



Source: Isola, Phillip, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros. "Image-to-image translation with conditional adversarial networks." CVPR, 2017.

Other Applications of Generative Models

■ Speech/Audio


- ◆ Oord et al. "Wavenet: A generative model for raw audio." arXiv, 2016.

■ Generating 3D from 2D

- ◆ Wu et al. "Learning shape priors for single-view 3d completion and reconstruction." ECCV, 2018.

■ Text-to-image

- ◆ Zhang et al. "Stackgan: Text to photo-realistic image synthesis with stacked generative adversarial networks." ICCV, 2017.



■ And many others...

- ◆ Scene completion
- ◆ Image editing
- ◆ Face aging
- ◆ Super-resolution
- ◆ Video prediction

Next Time

- Wednesday (4/10): Lecture
 - ◆ Topic: Diffusion Models

- Upcoming:
 - ◆ **Homework 5** due 4/12